

Distributed Power Architecture for a Multi-PSU Robotic System

Abstract

This report presents a conceptual overview of a distributed power architecture for a robotics prototype that uses multiple consumer-grade power supply units (PSUs) to power an integrated system. Instead of a single supply, several off-the-shelf PSUs are employed in parallel (by dividing the load, not by directly paralleling outputs) to drive three interconnected motherboard-based subsystems alongside motors, sensors, microcontrollers, and auxiliary devices. All PSUs share a common ground reference. We discuss the rationale for this multi-PSU approach in a prototyping context – including cost benefits and capacity scaling – and describe how power is distributed across subsystems. Key technical considerations are examined, such as the implications of using a common ground without isolation, potential interference and electrical noise between subsystems, and risks like backfeeding or instability in the absence of OR-ing diodes or relays. The challenges inherent to this architecture are analyzed, and we outline how the system could be refined in future iterations (for example, through improved grounding strategies, power management circuitry, or consolidation of supplies). The tone of the report is technical yet accessible, aiming to provide an academic-style perspective on leveraging multiple consumer PSUs in robotics prototyping. Citations are included to support the discussion in multiple reference styles (APA, MLA, IEEE) where appropriate.

Introduction

Robotic systems typically require multiple voltage rails and substantial currents to support components ranging from high-performance computing boards to motors and sensors. In professional designs, a **centralized power architecture** might use one high-capacity supply (or battery) feeding all components, or a **distributed power architecture** with a single high-voltage source and point-of-load converters at subsystems ¹. However, in early-stage **prototyping**, engineers often repurpose readily available **consumer-grade PSUs** (such as PC ATX power supplies) to meet power needs. This approach offers practical benefits: multiple smaller PSUs can provide combined current capacity across different voltages and can be cheaper or more accessible than a single purpose-built supply of equivalent power ² ³. For instance, one case study found that using two standard ATX PSUs doubled the available 5V current (from ~25–30A to ~50A) and spread the load and heat dissipation between units, all at roughly one-quarter the cost of a single high-wattage supply ² ³. In prototyping scenarios, leveraging spare PC PSUs “works in a pinch” even if it’s a bit of a kludge compared to a unified supply ⁴.

Rationale for Multiple PSUs: The decision to employ multiple PSUs in the *Huey* robotic shell (a complex system with three onboard motherboards and numerous peripherals) is driven by both practical and technical factors. First, no single consumer PSU on hand may meet the aggregate power requirements – especially if high currents are needed at standard voltages like 5V and 12V for motors, drives, and computing components. Using two or three smaller PSUs allows splitting of loads, effectively increasing total available current across rails ². This load-splitting approach has been observed in other high-power systems; for example, Backblaze’s storage pods used dual ATX PSUs not for redundancy but to **share the**

heavy load, with each PSU powering a subset of devices ⁵. Second, prototyping with off-the-shelf PSUs provides flexibility – one can mix-and-match units for different sections (e.g. one PSU dedicated to motors, another to computing boards) and easily swap or upgrade them as needed. Finally, there is a convenience factor: ATX PSUs offer multiple voltage outputs (+12V, +5V, +3.3V, etc.) with built-in regulation and protection, simplifying the power delivery to diverse robotic components during early development.

This report focuses on a conceptual understanding of such a multi-PSU power architecture. We will describe the overall system architecture and power distribution strategy, detail the common-ground scheme used to reference all subsystems, and examine the challenges that arise (including electrical noise coupling, interference between subsystems, and safety or stability concerns). We also discuss the limitations of this approach – for instance, the lack of isolation or diode OR-ing raises potential issues like backfeeding between supplies – and how these might be mitigated or designed out in future refinements of the system. The goal is to provide a clear technical overview suitable for an academic or engineering audience, without delving into specific implementation details of any one PSU model. In the following sections, we present the system architecture (§ **System Architecture**), analyze key challenges (§ **Challenges**), discuss the broader implications and possible improvements (§ **Discussion**), and finally summarize our conclusions (§ **Conclusion**).

System Architecture

Overview: The prototype robotic shell's power architecture consists of multiple consumer-grade PSUs powering different portions of the system in a **distributed manner**. Rather than tying all PSUs together on the same output, each PSU is allocated to specific subsystems or sets of components. In this design, three primary subsystems – each built around a motherboard (e.g., single-board computers or a server-grade motherboard and two auxiliary control boards) – are each powered by one or more dedicated PSUs. All PSUs share a **common ground** connection (0V reference) but their positive outputs are largely kept separate, each feeding its intended load group. This means, for example, PSU A might supply a 12V rail to the main motherboard and its peripherals, PSU B might supply another motherboard or high-power motors/actuators, and PSU C might power a third control unit and various sensors or auxiliary electronics. By partitioning the load in this way, the supplies are **not directly paralleled** on the same output node; instead, they operate in parallel at the system level by each handling a portion of the total load devices ⁶. This avoids the complexities of current-sharing between PSUs on a single rail and instead resembles a **decentralized power distribution**, where multiple sources feed different branches of the system.

Common Ground Strategy: A critical aspect of this architecture is that all subsystems must share a common reference potential for their signals to interoperate. The ground (negative) terminals of all PSUs are therefore **tied together**, creating a single common ground plane across the robot. This common ground ensures that any communication signals, sensor references, or control lines between devices powered by different PSUs have a mutual reference and can operate correctly. In practice, the PSU ground outputs are connected (often at a star point or common grounding bus in the chassis) so that 0V is uniform everywhere ⁶. It's worth noting that in ATX PSUs and similar, the DC ground is typically also connected to earth ground; since all PSUs are plugged into the same mains circuit, their earth grounds are inherently common. Nonetheless, explicitly tying their DC grounds together at a single point helps prevent differences in ground potential from arising across the robot's frame. This **common ground bus** approach is standard when using multiple supplies – one expert advises to “cross-tie the 0V ends” of all power supplies in such multi-PSU setups ⁶ to avoid floating references between them.

By not including any isolation or OR-ing diodes between the grounds, the design essentially treats all subsystem grounds as one unified node. Signals (e.g., a PWM control from a microcontroller on one board driving a motor driver on another supply) can flow reliably because the return currents all share that same ground path. The decision to avoid *diode OR-ing* or *relay isolation* on the outputs means there is **no intentional voltage drop or separation** between the supplies' outputs and the loads – every PSU is directly connected to its loads and to the common ground. This simplicity is part of the prototyping philosophy: fewer components and complexities make it easier to assemble and reconfigure the power system. It is akin to a rough implementation of a distributed power architecture: each PSU serves as a local power source for a cluster of components, somewhat analogous to point-of-load converters in a formal distributed architecture (albeit here the “point-of-load” source is a full PSU) ¹.

Power Distribution to Subsystems: Power wiring in this architecture is arranged so that high-current devices are served by the nearest or most capable PSU. For example, power-hungry actuators or drive motors (which run on 12V DC) might be wired directly to the 12V rail of the PSU that also powers a particular motherboard, thereby sharing that PSU's capacity. Other sensitive electronics like microcontrollers or sensors requiring 5V or 3.3V could either be fed from the same PSU's regulated outputs or from small DC-DC converters attached to that PSU's 12V rail. The three motherboards themselves each receive ATX power from their respective supplies. Importantly, because all grounds are common, **power distribution and signal distribution become intertwined**: a sensor powered by PSU A can safely send data to a processor on PSU B, as their grounds are common. The architecture relies on the fact that each PSU independently maintains its output voltage for its load; there is no centralized power management ensuring equal load sharing or failover. In essence, the robot has multiple “power islands” (one per PSU) that are connected at the ground level and via any inter-system signal wiring.

This kind of segmentation requires careful planning of which PSU feeds which component group. A guiding principle is to **avoid having two different PSUs drive the same exact load or node**. There are no parallel connections of two supplies to one circuit, which means classic current-sharing problems are averted. Instead, each device or module is clearly assigned to one PSU's domain. The common ground is the only intentional connection between the otherwise separate supply domains (aside from data connections). This approach has precedent in other multi-supply systems. For instance, in large disk storage arrays, one half of the drives can be run off one PSU and the other half off a second PSU – both PSUs' grounds are common, but each drive is powered by only one supply ⁵. Such a scheme ensures that **load distribution** is achieved without the supplies electrically fighting each other. It's recommended in engineering forums that if one must use multiple PSUs, “*split the load into two halves*” (or thirds, etc.) and dedicate each supply to a portion of the devices, rather than directly paralleling outputs ⁶.

No OR-ing Diodes or Redundancy: In this prototype configuration, **diode OR-ing** (i.e., using diodes or ideal diode controllers to combine outputs) is deliberately not used. In a redundant power supply design, OR-ing diodes allow one supply to take over if another fails, without backfeeding current into the failed unit ⁷. Here, because each PSU has its own dedicated outputs, redundancy was not a design goal – the aim was increased total capacity, not fault-tolerance. Avoiding OR-ing components saves cost and avoids the voltage drop that diodes introduce on the lines. It does mean, however, that **if one PSU were to lose power, the devices it feeds would shut down** and there is no automatic failover. The other PSUs would continue powering their sections, so the system as a whole might operate in a degraded state. For example, if the PSU powering the primary computer fails, the motors and other boards might still have power, but the “brain” is off – a potentially unstable scenario. Because of the common ground, a powered-down supply's output node will be at 0 V, effectively disconnected from its load (unless some other still-powered

connection unintentionally sources it current – a risk discussed later). In short, the architecture prioritizes simplicity and available hardware usage over advanced features like redundancy or hot-swap capability. Each PSU is a single point of failure for its portion of the system.

In summary, the system architecture can be viewed as a **distributed multi-source power network** inside the robot: multiple PSU “sources” feed various “loads” with all sources joined at ground. This yields a functional power solution for the prototype stage, leveraging the strengths of consumer PSUs (multiple voltages, high current on 5 V/12 V rails, built-in protections) while acknowledging their independent operation. Next, we explore the challenges and implications that arise from this design, particularly due to the common ground coupling and lack of isolation between subsystems.

Challenges

While the multi-PSU architecture successfully powers the robot prototype, it introduces several **electrical challenges and risks** that must be managed. These challenges stem from the interactions between the multiple supplies and the connected subsystems, especially via the common ground. Key issues include ground potential differences and loops, interference and noise coupling between subsystems, the risk of backfeeding currents between PSUs, stability during load transients, and difficulties in power monitoring and control. We discuss each of these in turn.

1. Ground Potential Differences and Ground Loops: With all PSUs sharing a common ground, it is crucial to ensure that this ground is truly at the same potential throughout the system. In practice, if the ground wiring is not ideal, high current flows can cause slight voltage differences between different ground points – a phenomenon known as ground offset or ground bounce. For instance, when a heavy motor current returns to its PSU, the finite resistance and inductance of the ground wire means that section of ground might rise a few millivolts (or more) above the rest of the ground plane. If at the same time a microcontroller ground at a distant point is “quiet,” the difference in ground potential can disturb the reference for signals between them ⁸. One answer on Electrical Engineering Stack Exchange explains this using Ohm’s law: *any shared ground conductor can be modeled as a small resistor, so large currents through it cause a voltage drop that appears as noise to other devices on the same ground* ⁹. The effect is analogous to lights dimming in a house when a heavy appliance turns on – the wiring impedance causes a temporary drop in the shared supply. In a robot, a sudden motor stall current could momentarily raise the local ground, causing sensor readings or logic levels to shift unpredictably.

Moreover, with multiple PSUs, there is a possibility of **ground loops** if the ground is connected in more than one topology. Ideally, all grounds should connect at a single common point (a star ground configuration) to avoid loop currents. If, for example, PSU A and PSU B are both bolted to the metal chassis which itself is used as ground, and additionally a cable connects their grounds, a loop is formed. AC currents or noise could circulate in this loop. The prototype uses a common ground bus to minimize such loops, but care must be taken in routing ground wires. A recommended practice is star grounding, where each PSU ground and each major return path meets at one node (or bar) on the chassis ¹⁰. This limits the ground current paths such that high return currents (from motors, for example) do not flow through sensitive ground reference points for electronics. Without isolation, **conducted interference** is a concern – one PSU’s noisy ground currents (from a switching regulator or a motor driver it powers) can introduce noise into another PSU’s ground reference. Indeed, engineers caution that tying grounds together is usually fine, but “*be cognizant of what return currents are going where and by what path*”, especially with high dI/dt loads like

motors ¹¹ . If poorly managed, these currents can cause different ground nodes to momentarily differ in voltage, upsetting the system's stability.

Mitigation strategies for ground issues include using sufficiently thick, low-resistance ground wiring (or ground planes) to carry heavy currents, thus reducing voltage drops. The prototype should also physically cluster high-current devices closer to their PSU ground connection to shorten return paths. Star grounding or a single robust ground bus bar is employed to collect all returns. Additionally, separating analog/sensor ground returns from power ground returns until a single combine point can help keep noise away from sensitive circuits (this is a form of ground isolation without an actual isolation barrier). In summary, maintaining a solid common ground in a multi-PSU system is feasible but requires careful attention to wiring layout to avoid ground loops and **common impedance coupling** of noise ⁹ .

2. Power Interference and Electrical Noise: A direct consequence of the common ground and the lack of isolation is that electrical noise generated in one part of the system can propagate into others. Each PSU is a switching power supply (in the case of ATX units), which inherently generates some switching noise on its outputs. When multiple such supplies and various loads (especially motors, which generate electromagnetic interference when their brushes commutate or their drivers chop currents) are all connected, the noise can sum and transfer through the common connections. For example, if a DC motor (powered by PSU B) is rapidly changing speed, the current surges and voltage spikes from that motor might inject noise onto PSU B's 12V line and ground. Through the common ground, a portion of this noise could appear as ground voltage fluctuations that superimpose on PSU A's outputs as well, potentially disturbing the electronics powered by PSU A. Similarly, any high-frequency ripple or transient from one PSU can couple into other subsystems via the shared ground or any other wiring commonalities (like if subsystems share a signal line, the reference noise can effectively appear as signal noise).

Interference can also occur via radiated paths – multiple PSUs with their own fans, transformers, and switching circuits can create a complex EMI environment inside the robot's chassis. Without careful cable routing or filtering, sensitive sensor lines could pick up this electromagnetic noise. The architecture currently does not use isolation transformers or optocouplers between subsystems (since everything is grounded together), meaning *all electronics see each other's noise*. This is an acknowledged trade-off of the simple approach: keeping circuits separate would eliminate conducted interference, but “unless there's a technical reason to have a common ground, it's almost always best to keep circuits separate” one engineer notes – in our case we **do** have a reason (intercommunication) to tie grounds, so we accept the coupling and try to manage it ¹² .

To manage noise, the prototype can incorporate **decoupling capacitors, filters, and shielding**. Each PSU output feeding a load should have local capacitors to absorb transients (e.g., electrolytic and ceramic capacitors near motor drivers, etc.). Ferrite beads or inductors can be used on supply lines to choke high-frequency noise. In more advanced iterations, one could introduce **isolated DC/DC converters** for particularly noise-sensitive components (for instance, running a sensitive sensor array from an isolated regulator fed by one PSU – the sensor would still ultimately reference ground for data lines, but its power noise would be isolated). Additionally, good practices like twisting wire pairs and using shielded cables for signal lines can reduce noise coupling. In the current conceptual design, these measures are noted as necessary refinements: at the prototype stage, some increased noise floor might be tolerated, but it underscores the need to refine power distribution in the future to ensure signal integrity.

3. Backfeeding and Cross-Regulation Risks: **Backfeeding** refers to current flowing from one power source into another, which is a major risk when multiple supplies are connected without OR-ing diodes or isolation. In our architecture, we intentionally avoid directly paralleling any two PSUs on the same supply rail. This greatly reduces the chance of backfeeding under normal operation, since each PSU “sees” only the loads assigned to it. However, subtle scenarios still exist where backfeed could occur. One scenario is if one PSU is powered off (or fails off) while another PSU is still on, and there are connections through the loads or signal lines that provide a path. For example, consider a microcontroller board powered by PSU A that has a GPIO line connected to a motor driver powered by PSU B. If PSU A is off, the microcontroller is unpowered – but the motor driver on PSU B might drive that GPIO line high. In doing so, it could inadvertently source current through the protection diodes of the microcontroller’s input pin, effectively **feeding 5 V back into PSU A’s 5 V rail via the microcontroller**. This kind of indirect coupling can power parts of a supposedly off subsystem at a phantom low level (often called parasitic or “ghost” powering). It’s not the catastrophic type of backfeed that drives large currents, but it can lead to unpredictable behavior (e.g., an MCU partially powering through I/O, which can damage components).

The more dangerous form of backfeeding would be if two PSUs were accidentally tied together on a high-current line (say two 12 V lines joined). Without diodes, the one set at slightly higher voltage would push current into the other. ATX PSUs are not generally designed to sink current on their outputs; if forced, it could **damage the supply or trigger protective shutdowns** ⁷. A powered-off PSU typically presents near 0 V at its output, so if another live PSU is connected, it will drive current into the off PSU’s output capacitors and circuitry. As an expert succinctly put it, *“you can’t back-feed voltage/current into a powered-off supply... or you can, but the power supply may get destroyed.”* ⁷. Therefore, in our design, we absolutely avoid wiring two PSU outputs together at different potentials. The common ground is the only shared line; all positive rails go to distinct loads.

However, even with these precautions, one must consider failure modes: if PSU A were to develop a fault and output an abnormal voltage, could it feed into devices that connect to PSU B and then into PSU B? This is unlikely unless there is a direct device bridging (which our load partitioning prevents), but it highlights why in more mature designs OR-ing diodes or FET-based ideal diodes are used – they block reverse currents entirely. Our prototype takes the calculated risk of running without those, relying on careful wiring separation. The consequence is that **there is no protection if a PSU outputs out-of-spec voltages or if someone accidentally cross-connects outputs**. For the prototype stage, the team is presumably managing this by design discipline (clear labeling, harness design) and accepting the risk in exchange for simplicity.

4. Load Transient Handling and Stability: With multiple independent PSUs, system-wide power stability can be an issue under fast load changes. Each PSU regulates its own outputs, but consider a scenario where a large motor on PSU B quickly draws a surge of current (for instance, on startup or when hitting a mechanical stop). PSU B’s voltage might dip momentarily under the sudden load (even if within spec, there’s transient droop). PSU A, which might be powering the main computer, is unaffected in terms of its own output voltage – however, the motor surge might have caused a ground bounce as described earlier, or even a slight drop in mains voltage that affects all supplies. The result could be the main computer experiencing a brown-out or a sensor giving erratic readings at the exact moment of the motor surge. In a single-supply system, all loads share the same transient pool and the supply can be designed to handle it (or brown-out the whole system uniformly). In a multi-PSU system, one part of the system browns out while others remain up, which can be a worse failure mode (partial operation, possibly causing uncontrolled behavior).

Additionally, the **power-up sequencing** can be a challenge: when turning on the system, if the PSUs are switched on at slightly different times, some subsystems will come up before others. For instance, one PSU might ramp to 12 V faster, powering its motherboard, while another PSU is a second behind. This could lead to the first subsystem sending signals to the second before the second is powered, again risking backfeeding through signal lines or at least causing communication errors. In our prototype, without a coordinated power-on mechanism (no relay or single master switch that guarantees simultaneous start), the user must take care to turn on all PSUs together. There are products (like “Add2PSU” boards in PC modding) that link PS_ON signals so that a primary PSU can auto-start a secondary – but the design explicitly did not incorporate such relay triggers. Thus, an operator turning on the robot must remember the correct procedure to energize all supplies. During operation, if one PSU were to shut down (due to overload or overheat), there is no graceful handover of its load; that part of the system simply loses power. This lack of **power redundancy** or sharing means the architecture is inherently less stable in the face of individual PSU faults.

To improve stability, one might overspec each PSU (use only, say, 50% of each PSU’s max load in normal operation) so that transient demands are easily met without significant voltage sag. Also, distributing the most variable loads (like motors) on separate PSUs from sensitive loads (like computing) can localize disturbances. This seems to have been done by dedicating different PSUs to different functions. If motors and logic must run on one PSU, adding large capacitors or even supercapacitors on the motor supply can buffer the surges. In terms of start-up sequencing, future revisions could employ a master switch or a small microcontroller circuit that activates all PSUs synchronously (by shorting their PS_ON to ground at the same time). For now, the prototype likely uses a manual or improvised method (e.g., jumping PS_ON pins and using a power strip to turn on AC to all supplies simultaneously). These are workable but not ideal solutions, underscoring that **power management is rudimentary** in the current setup.

5. Monitoring and Diagnostics: A subtle challenge is the **lack of unified power monitoring** in a multi-PSU scheme. Each PSU has its own feedback and protections (shutdown on overcurrent, etc.), but there’s no single dashboard telling the robot operator how power is distributed or if a PSU is nearing its limit. For instance, one PSU could be running near maximum load (perhaps the one powering many motors), while another is lightly loaded. Without monitoring, this imbalance might go unnoticed until the overloaded PSU triggers a shutdown. In a one-supply system, it’s easier to instrument the current draw and voltage of that one supply. In the multi-PSU system, one would need multiple sensors – one per PSU output – to get a full picture. The prototype likely doesn’t include an elaborate power monitoring network (given the focus on basic functionality). This introduces the risk that an overload condition on one supply is not detected in time. It also complicates energy management: if the robot were to eventually run on battery power, knowing how each subsystem draws power is important for efficiency, but with separate PSUs, the data is siloed unless explicitly gathered.

Furthermore, if something goes wrong – say a short circuit in a sensor board – the corresponding PSU will shut down to protect itself. The rest of the system stays on, but now part of the robot is unpowered. This partial failure mode can be hard to troubleshoot if you don’t have indicators for each PSU’s status. It may not be immediately obvious why a certain subsystem is unresponsive (is it its PSU tripped, or the board itself died?). Including LED indicators or voltage readouts for each PSU’s output in the prototype can alleviate this. In a refined design, a supervisory microcontroller could monitor all supply voltages and perhaps communicate to the main system if any rail falls out of range.

Summary of Challenges: In essence, the multi-PSU architecture trades the convenience of using available supplies for a set of challenges in electrical integration. **Grounding and noise** issues require careful mitigation to prevent interference between subsystems. The absence of **isolation/OR-ing** means careful wiring discipline is needed to avoid backfeed and to handle failures gracefully (which, in truth, they won't be handled gracefully – it's a known limitation). **Stability** under dynamic loads is not as robust as it could be, and there is no inherent redundancy or load-sharing beyond static assignment. Finally, **monitoring** multiple independent sources is more complex. These challenges do not mean the architecture is unworkable – only that it must be acknowledged as a prototype solution. Many of these issues can be managed with good engineering practices (as noted, e.g., heavy ground conductors, star topology, decoupling capacitors, prudent load distribution). In the next section, we discuss the implications of these challenges and how the architecture could evolve to address them in future iterations.

Discussion

The use of multiple consumer-grade PSUs in a single robotic system highlights an interesting balance between short-term pragmatism and long-term design ideals. From a **conceptual standpoint**, this architecture embodies a decentralized power approach: instead of one central supply, power conversion is distributed across several units located near the loads they serve. In fact, it echoes some principles of formal Distributed Power Architecture (DPA) used in high-end electronics, albeit in a cruder form. In a DPA, a single high-voltage source feeds many local converters at the loads ¹. Here, we effectively have multiple “local” sources (the PSUs) each directly powering certain loads. The benefits realized are immediate – ease of using existing hardware, increased total capacity by splitting loads, and reduced need for custom power electronics at the prototype stage. The rationale discussed earlier for prototyping (cost, availability, flexibility) has indeed played out in practice: the team achieved a working power solution without waiting for or designing a specialty power supply.

Why It Makes Sense in Prototyping: Prototyping often involves rapidly implementing a solution with components on hand, then iterating. Consumer PSUs are **modular and well-understood** – one can add a new PSU to the system to get an extra 12 V rail in a matter of minutes. This modularity is helpful when the power budget of the robot is in flux (as features and hardware are added, power needs grow). Using multiple smaller PSUs also sidesteps the need to immediately invest in a costly high-power supply. As noted earlier, two standard PSUs can sometimes outperform a single larger unit in specific ways (more 5 V current, more connectors, etc.) ². Additionally, the **common ground scheme** makes it straightforward to integrate subsystems: there is no need for complex level shifting or isolation circuits for every signal, because everything shares reference ground. This simplicity accelerates development – you can wire components together freely knowing their grounds are common, which is often assumed in typical development setups.

However, as the previous section detailed, these conveniences come at the cost of potential reliability and performance issues. **Electrically, the system is less elegant and less robust** than it could be. During prototyping, one might accept occasional quirks (like a noise-induced sensor glitch when a motor starts) as long as the overall system works. But for a production or advanced prototype, those would need ironing out.

Evolution and Refinement Opportunities: Looking ahead, there are several pathways to evolve this power architecture for greater robustness:

- **Consolidation into a Single Supply or Unified Bus:** One straightforward improvement is to replace multiple PSUs with a single higher-capacity supply that can handle the entire load. By doing so, all components naturally share the same power source, eliminating any ground differential issues (beyond internal wiring drops) and removing the chance of inter-supply backfeed. An expert on an electronics forum suggested exactly this when faced with a multi-PSU rack system, noting that a single 750 W–850 W ATX supply can provide 60–70 A on 12 V and would “solve the grounding issue” while simplifying replacement and reducing complexity ¹³. For the robotic shell, if a suitably powerful supply (or a custom-designed power distribution unit) is obtained, it could feed all three motherboards and peripherals from one source. The downside is the initial reason multiple PSUs were used – maybe no single supply was available or one supply might not provide different voltage rails needed. A compromise is to have one central high-power DC source (say 48 V battery or 24 V supply) and then use DC-DC converters at each subsystem (basically implementing a proper distributed architecture with point-of-load regulators). This would provide isolation between subsystems in terms of noise (to some degree, depending on converter isolation) and simplify grounding (only one ground reference from the single source). The **intermediate bus architecture** approach, common in industry, could be a model: use an intermediate voltage and multiple converters ¹⁴. Adopting this in the robot would be a significant redesign, but conceptually it’s the direction to move once prototyping transitions to productization.
- **Incorporating OR-ing and Redundancy:** If there is a desire to keep multiple PSUs for either **redundancy or increased capacity**, then adding OR-ing diodes or ideal diode controllers is a logical next step. For instance, two 12 V PSUs could be diode-ORed to feed a common 12 V bus – this would allow one PSU to take over if the other fails, improving reliability. It also prevents backfeed because the diodes block reverse currents ⁷. The trade-off is the voltage drop and some efficiency loss across the diodes (Schottky diodes drop ~0.3–0.5 V at high current, which can be significant). Modern ideal diode circuits use MOSFETs to reduce this drop. Implementing this would require a bit of custom circuitry or off-the-shelf OR-ing boards. Given this is a prototype, full redundancy might not be a priority now, but if the robot’s design moves toward any kind of **hot-swappable power or battery backup**, OR-ing is essential. For example, integrating a battery as a backup would involve OR-ing the battery with the mains supply so that one can seamlessly take over from the other. The current architecture has no provisions for that, but it could evolve in that direction.
- **Improved Grounding and Isolation:** To tackle interference issues, future designs can segment grounds into analog and digital, or dirty (motor) and clean (logic) grounds, that meet at a single point. Right now everything is tied together; introducing *ground isolation techniques* could help. While optical or galvanic isolation of every signal is impractical (and contrary to having a common ground), selective isolation might be used. For instance, isolating the motor driver control signals via optocouplers would allow one to physically separate the motor ground from the logic ground, eliminating a huge noise source. If that’s too drastic, at least ensuring a star grounding scheme (if not already in place) should be formalized in the design. Additionally, using differential signaling for any long data lines (so that they are less sensitive to ground noise) would be beneficial.
- **Active Power Management:** The prototype currently likely relies on manual control of PSUs. A more advanced iteration could include a **power management board** that interfaces with all PSUs. This

board could do several things: (a) turn on/off all PSUs in a controlled sequence, possibly with a staggered start to avoid surges, (b) monitor voltages and currents from each PSU, (c) provide alerts or even take action (like shedding non-critical loads) if a PSU is stressed, and (d) balance the load if possible (for example, if one PSU is near limit and another has headroom, perhaps some devices that can be powered by either could be switched over – though in our architecture devices are fixed to one PSU, so this would require hardware changes such as diodes or switches to allow power source switching). While implementing full load balancing is complex, monitoring is very feasible. Many modern systems use something like an I2C-based power monitor chip per rail. The data could be logged by the robot's main controller to analyze power usage patterns.

- **Addressing Thermal and Mechanical Integration:** An often overlooked aspect of multiple PSUs is the physical space and heat. Three PSUs will take more space and likely produce more waste heat (with lower overall efficiency) than a single supply sized for the job. As the design iterates, one might consider using *server power supply modules* which are compact and efficient. These often support current sharing out-of-the-box and are designed for parallel operation (for example, two server PSUs each capable of full load, in an N+1 redundant config). Replacing consumer ATX units with industrial or server supplies could both save space and add reliability. The trade-off is that those might need custom breakout boards and lack the multiple voltage rails (often server PSUs just give 12 V, requiring separate DC-DC converters for 5 V, etc.). In the long run, a custom power distribution unit could be built, but in the interim, choosing higher-quality PSUs intended for parallel use is an improvement path.
- **Safety Enhancements:** The current approach, lacking isolation, means that a fault in one part could potentially propagate. For instance, if a high-voltage spike (from a motor driver fault) occurs on one PSU's output, it could travel through common ground or signal lines to other subsystems. To improve safety, transient voltage suppressors (TVS diodes) could be placed on supply lines to clamp spikes, and fuses or circuit breakers can isolate faults. Each PSU's output to its loads should ideally have fusing so that a short in one branch doesn't drag the PSU (or others) down. It's likely the prototype uses the PSUs' internal protection and maybe some fuses, but as currents and complexity grow, a coordinated protection scheme is valuable. For example, a polyfuse on the 5 V line to a sensor cluster can prevent one bad sensor board from shorting out that 5 V line and pulling down the whole PSU.

Academic Perspective: Conceptually, this multi-PSU architecture is a case study in leveraging **distributed, heterogeneous power sources** in an integrated system. It invites comparison with classical power distribution topologies. We can analyze it through the lens of system engineering: it achieves *functional power delivery* and some degree of *modularity*, at the expense of violating some design rules about isolation and coupling. In an academic context, one could ask: what does this approach mean for system reliability and performance? We have identified that it creates multiple single points of failure (each PSU is independent), whereas a proper redundant design might eliminate single points of failure but at cost of complexity. It also highlights a **scaling issue**: this approach might work for 2–3 PSUs, but scaling to many more would become unmanageable in terms of ground wiring and coordination. Thus, one insight is that while using multiple small PSUs is expedient early on, there is a crossover point where designing a consolidated power system becomes necessary.

Comparison to Alternatives: If we consider alternatives – e.g., one large PSU, or batteries plus regulators – the multi-PSU approach is perhaps less optimal for final performance, but possibly *more adaptable* during

development. Teams can swap out one PSU for another easily (maybe upgrading one 500 W PSU to 650 W if needed, independently of the others). It also allows powering subsets of the system individually (for testing, you could turn on just one PSU to power one part while keeping others off). This granularity can be useful in debugging and bringing up a complex system incrementally. In contrast, a one-supply system is all-or-nothing (though you could still switch sub-circuits off individually if designed that way).

In the future, as the project moves beyond prototyping, we expect a convergence towards a more **integrated power architecture**. This could mean designing a custom power distribution backplane that all components plug into, possibly powered by one or two sources. That backplane could handle ground isolation sections, filtering, and provide connectors for each subsystem, simplifying wiring (right now, with multiple ATX PSUs, the wiring is likely a mess of cables). The current approach has been invaluable for learning what the power demands are and how the system behaves under load. Those data will inform the requirements for any subsequent power system design.

Conclusion

In this report, we examined the conceptual design of a distributed power architecture employing multiple consumer-grade PSUs in a complex robotic shell prototype. This multi-PSU approach was chosen for pragmatic reasons: it allowed the use of readily available hardware to meet the system's considerable and diverse power needs, and enabled partitioning the load across several supplies. We detailed how the architecture is arranged – with three PSUs supplying three motherboard-centered subsystems plus motors and devices, all tied together with a common ground. This design achieves the basic goal of powering the robot's components and has parallels to distributed power strategies seen elsewhere in computing and robotics prototypes (for example, splitting loads between PSUs as done in certain server pods ⁵). The use of a common ground simplifies subsystem integration but also emerges as a focal point for many challenges.

Our exploration of the **challenges** highlighted important considerations. The common ground, while necessary, can lead to ground loops and noise coupling, especially with high-current motor loads causing ground bounce ¹¹ ⁸. The lack of isolation or diode OR-ing means there is a very real risk of backfeed damage if supplies are not perfectly kept to their own loads ⁷. Stability issues can arise since the supplies operate independently – a surge on one rail can indirectly perturb another subsystem. Monitoring multiple supplies is more complicated, potentially leaving blind spots in power management. These factors underscore that the architecture, as implemented, is a **compromise solution** characteristic of a prototype: it favors immediate functionality and flexibility over refined performance and fault tolerance. In academic terms, it sacrifices some system robustness (violating the ideal of a single coordinated power source) in exchange for modularity and convenience.

We also discussed how this architecture could be **iterated upon and improved**. Future refinements might include consolidating power into a single source or a more formalized bus distribution to eliminate inter-supply issues. Alternatively, if multiple PSUs remain, adding proper OR-ing circuits and control logic would make the system more resilient and safe. Improved grounding practices (like star topology) and filtering will be critical to manage interference as the system scales. Essentially, as the project matures, the power system would need to evolve from an ad-hoc collection of PC supplies into a purpose-designed power distribution network befitting a complex robotic platform. This could involve custom hardware but would yield benefits in reliability and performance. Recommendations from experienced engineers – such as using

a higher-capacity single PSU to “make it easy” and avoid multi-ground complications ¹³ – should be weighed against the project’s goals for redundancy and modularity.

In conclusion, using multiple consumer PSUs with a common ground has proven to be a **workable strategy in the prototype stage**, demonstrating the viability of quickly mobilizing available resources to power a multifaceted robotic system. The approach is conceptually simple and provides a learning platform to identify power requirements and interactions in the system. However, it comes with clear limitations in terms of electrical elegance and safety. As this distributed power architecture is refined, the lessons learned – about load distribution, interference management, and failure modes – will guide the design toward a more robust solution. The experience underscores a common engineering theme: the solutions that get a prototype up and running (multiple cheap PSUs tied together) are often not the same solutions one will use in the final implementation, yet they are a critical step in the development process. This report provided a detailed overview of the concept, supported by technical insights from literature and practice, in an effort to bridge that understanding towards eventual implementation of a safer, more stable power architecture for the robotic system.

References:

- Stilez. (2018). *How safe is parallel wiring two ATX PSUs...?* Electrical Engineering Stack Exchange ² ³ . (APA)
- Andy aka. *Electrical Engineering Stack Exchange answer*. 2018. *Parallel vs Split Loads with Dual PSUs* ⁶ . (MLA)
- Justme, *StackExchange user*. (2024). *You need two diodes because you can't back-feed...* ⁷ . (IEEE)

¹ ¹⁴ sakai.mci4me.at

https://sakai.mci4me.at/access/lessonbuilder/item/1183751/group/special.project.paper_mech/Graduates%20of%202024/TemsamaniPhillip_MA-paper.pdf

² ³ ⁵ ⁶ power supply - How safe is parallel wiring two ATX PSUs, if they weren't explicitly sold as dual PSUs? - Electrical Engineering Stack Exchange

<https://electronics.stackexchange.com/questions/351600/how-safe-is-parallel-wiring-two-atx-psus-if-they-werent-explicitly-sold-as-dua>

⁴ Dual computer PSUs, switching, and common ground - Page 1

<https://www.eevblog.com/forum/projects/dual-computer-psus-switching-and-common-ground/>

⁷ Why are 2 diodes needed instead of 1 when diode-ORing two power supplies? - Electrical Engineering Stack Exchange

<https://electronics.stackexchange.com/questions/711680/why-are-2-diodes-needed-instead-of-1-when-diode-oring-two-power-supplies>

⁸ ⁹ ¹⁰ ¹¹ ¹² arduino - Can 3 power supplies safely share common ground? - Electrical Engineering Stack Exchange

<https://electronics.stackexchange.com/questions/716154/can-3-power-supplies-safely-share-common-ground>

¹³ Grounding in a multiple power supply system - Page 1

<https://www.eevblog.com/forum/beginners/grounding-in-a-multiple-power-supply-system/>