

Real-Time Audio Understanding and Generation in AI: A Comparative Study of OpenAI, Google Gemini, and DeepSeek

Abstract

Artificial intelligence has made significant strides in audio processing, enabling machines to **understand** and **generate** audio with unprecedented accuracy and realism. This paper presents a comprehensive overview of state-of-the-art AI models for audio, focusing on *real-time* speech understanding (e.g. speech recognition, speaker identification, sound classification) and audio generation (e.g. text-to-speech, voice cloning, music synthesis). We emphasize OpenAI's contributions – notably the Whisper speech recognition model and the Voice Engine for voice cloning – and compare them against contemporaneous developments from Google's **Gemini** platform and China's **DeepSeek** models. We examine the technical architectures of these systems, their performance, latency, data efficiency, and scalability, as well as their offline and real-time capabilities. Real-world applications are discussed, from conversational voice assistants and transcription services to creative tools for music generation. We also address ethical considerations surrounding synthetic audio (such as voice cloning and AI-generated music) and outline current limitations and future directions. The comparative analysis highlights how OpenAI, Google (Gemini), and DeepSeek each approach the audio domain: OpenAI with specialized models and cautious deployment, Google with integrated multimodal systems and extensive data, and DeepSeek with massive mixture-of-experts models. **Ethical safeguards** like watermarking of AI audio and usage policies are increasingly crucial in parallel with technical advances. The paper concludes by underscoring the balance between innovation in audio AI and the responsible use of these powerful technologies.

Introduction

Audio is a rich and vital modality for human communication and expression. Advances in artificial intelligence have allowed machines to *listen*, *understand*, and *generate* audio in ways that approach human-like performance. From virtual assistants that can transcribe and respond to speech in real-time, to algorithms that synthesize natural-sounding speech or even compose music, AI-driven audio technologies are rapidly evolving. In recent years, deep learning architectures (especially large-scale Transformer models) and vast training datasets have led to breakthroughs in speech recognition and generation. These developments open up new possibilities for human-computer interaction, accessibility, content creation, and more, but also raise new technical challenges and ethical concerns.

OpenAI has been at the forefront of this audio AI revolution. In 2022, OpenAI introduced **Whisper**, a neural network for automatic speech recognition (ASR) trained on 680,000 hours of multilingual audio data. Whisper demonstrated robust transcription performance across many languages and acoustic conditions, showcasing the power of large-scale supervised training. More recently, OpenAI developed a powerful text-to-speech model called **Voice Engine**, capable of cloning a speaker's voice from a few seconds of audio. Voice Engine's technology powers ChatGPT's voice feature and can generate speech nearly indistinguishable

from the original speaker after hearing only a 15-second sample. While these tools push the state of the art, OpenAI has taken a cautious approach to their deployment – for example, limiting the release of Voice Engine due to potential misuse.

Meanwhile, other organizations have advanced their own audio AI systems. Google's latest **Gemini** models (developed by the Google DeepMind team) are designed to be *natively multimodal*, handling text, images, **audio**, and even video within a unified model. Announced in 2023 and updated through 2025, Gemini integrates Google's long experience in speech technologies (such as the **Universal Speech Model** for ASR and high-fidelity TTS systems) into a general-purpose AI. Gemini 2.5, for instance, can engage in real-time voice conversations, understand spoken queries, and generate expressive speech responses with very low latency. In parallel, a Chinese AI company called **DeepSeek** has emerged with enormous-scale models that also encompass audio capabilities. DeepSeek's flagship model, DeepSeek-V3, is a mixture-of-experts language model totaling 671 billion parameters (with ~37B parameters active per token). DeepSeek's models are multimodal and have demonstrated notable achievements in AI music generation – a recent version garnered attention for producing music that rivals human compositions, signaling a leap forward in AI-driven music creation ¹.

This paper provides an academic-style exploration of these developments in AI audio understanding and generation. We begin by reviewing the technical architectures of leading models, focusing on how they achieve tasks like transcription, speaker identification, and sound classification (audio understanding) and tasks like text-to-speech, voice cloning, and music synthesis (audio generation). We then present a comparative analysis of OpenAI's, Google's (Gemini), and DeepSeek's approaches along key dimensions: model performance/accuracy, real-time latency, data efficiency, and scalability. The offline versus real-time capabilities of each platform are discussed in detail, since these determine practical deployment in applications like live conversation or on-device processing. Next, we survey real-world applications enabled by these audio AI systems – from voice assistants and transcription services to creative and assistive technologies. Finally, we address the ethical considerations (e.g. the risk of deepfake voices, privacy and copyright issues) and limitations of current models, and outline future directions for research and responsible innovation in this domain. By comparing the strategies of OpenAI, Gemini, and DeepSeek, we aim to illuminate the state of the art in AI audio and the trade-offs involved in pushing towards more **real-time**, high-quality, and safe audio AI systems.

Technical Background: AI Models for Audio Understanding

Speech Recognition and Audio Understanding Architectures

Automatic Speech Recognition (ASR) – converting spoken audio into text – is a foundational task in audio understanding. Modern ASR systems predominantly use neural network architectures trained on large datasets of audio and transcripts. OpenAI's Whisper is a representative state-of-the-art ASR model. Whisper employs a *sequence-to-sequence Transformer* architecture with an encoder-decoder design. Input audio is first converted to a log-Mel spectrogram, which the encoder processes; the decoder then generates text tokens (characters or word pieces) corresponding to the transcribed speech. Uniquely, Whisper was trained on a massive **680,000 hours** of supervised data from the web, spanning multiple languages and tasks. This vast and diverse training set gave Whisper strong robustness to accents, background noise, and technical jargon, and even enabled it to perform multilingual transcription and translation within a single model. Because Whisper was not fine-tuned to specific datasets, it does not top specialized benchmarks like LibriSpeech for English, but it achieves much more robust **zero-shot** results across varied real-world data

(reportedly ~50% fewer errors on diverse test sets compared to prior ASR models). Whisper's architecture processes audio in 30-second chunks, which has implications for real-time use (discussed later). The open-source release of Whisper (in models ranging from 39 million to 1.5 billion parameters) has allowed developers worldwide to build speech-enabled applications on-device, marking a significant milestone in accessible ASR technology.

Google's approach to speech recognition, recently unified under its **Universal Speech Model (USM)**, shares similarities but also key differences. USM is a family of encoder-decoder models (around 2 billion parameters each) capable of transcribing over 100 languages. Google leveraged *self-supervised learning* on 12 million hours of unlabeled speech to pre-train USM's Conformer-based encoder, then fine-tuned on a smaller amount of labeled data (28 billion text sentences and a few thousand hours of transcribed speech per language). This two-stage training (unsupervised audio pre-training followed by supervised fine-tuning) enabled **coverage of 300+ languages**, including many low-resource languages. The Conformer encoder architecture (a convolution-augmented Transformer) allows USM to handle long audio input efficiently and even produce outputs via different decoders (CTC, RNN-T, or LAS) for streaming or offline modes. The result of this training strategy is a model that not only generalizes to new languages with minimal data, but also achieves state-of-the-art accuracy. On YouTube caption data, USM obtained under 30% word error rate on average across 73 languages – a milestone for multilingual ASR. Notably, when directly compared on overlapping languages, Google reported USM had a **32.7% lower WER** than OpenAI's Whisper (large-v2) on those languages, demonstrating the gains from USM's massive unlabeled training and fine-tuning approach. In practice, Google has deployed these models in services like YouTube captioning and Cloud Speech-to-Text. The **Gemini** AI model incorporates such speech recognition capability as one modality of its multimodal system. Gemini 2.5 is described as understanding streaming audio input and even combining it with video input for richer context. Unlike Whisper's fixed 30-second segmentation, Google's production ASR (e.g. the *Chirp* model, an evolution of USM) supports continuous *streaming recognition*, outputting transcripts in real-time with low latency for interactive applications ². Additionally, Google's ASR services include features like **speaker diarization** (assigning transcript segments to different speakers) and **keyword spotting**, though these are often handled by auxiliary models or pipelines alongside the core ASR engine.

DeepSeek's strategy for audio understanding is somewhat different, reflecting its origin as a large language model with added multimodal capabilities. DeepSeek-V3 is fundamentally a gigantic **mixture-of-experts (MoE) Transformer** aimed at general AI tasks. It can process various data types – text, images, and audio – by converting them into a machine-readable form and feeding them into the model's layers. In practice, DeepSeek's audio understanding likely relies on converting audio to text (via ASR) as an intermediate step for most language-heavy tasks like transcription or summarization. Indeed, developer guides illustrate using DeepSeek-R1 in combination with an external streaming transcription API (like AssemblyAI) to build voice agents, implying that DeepSeek may *not* yet have a dedicated built-in ASR as efficient as Whisper or USM for raw audio. However, DeepSeek's multimodal design suggests it can **ingest audio features** directly for analysis if needed. The model's architecture presumably includes an audio encoder (potentially similar to a Transformer or conformer) feeding into the same latent space as text, enabling tasks such as transcribing audio, answering questions about audio content, or classifying sounds within an audio stream. DeepSeek documentation highlights "multimodal understanding" – for example, the ability to transcribe audio files and then summarize or answer questions about them in one workflow. This indicates that DeepSeek treats audio understanding as one facet of its general problem-solving, using its massive parameter count to reason about audio-derived information. One consequence is that DeepSeek's core model (especially V3 with MoE) is **extremely large**, which could impede real-time deployment for audio understanding unless

distilled versions are used. DeepSeek-R1-Distill, a compressed variant released in late 2024, was made available for local deployment, suggesting that offline use is possible with a downsized model, though possibly at some cost to accuracy. In summary, while OpenAI and Google offer specialized, highly-optimized ASR solutions, DeepSeek's approach integrates audio understanding into a broader AI model, leveraging scale (671B parameters total) over specialization. This may afford DeepSeek flexibility (e.g., understanding audio in context of other modalities), but it likely relies on the foundational advances from models like Whisper or USM to achieve raw transcription capability.

Beyond speech transcription, audio **classification** and **speaker identification** are also important tasks in audio understanding. These involve recognizing *what* is present in an audio scene (e.g., classifying sounds like music, applause, sirens, or detecting keywords) and *who* is speaking in a recording. Traditional approaches for these tasks use learned audio embeddings and classifiers (for example, Google's **AudioSet** and **YAMNet** models classify environmental sounds, and speaker ID often uses embedding networks like x-vectors). In the context of our focus: OpenAI's Whisper produces segment-level language ID and can output timestamps, but it does not explicitly perform speaker identification by itself (users must apply separate diarization algorithms to the transcript). Google's ecosystem does incorporate speaker recognition: as noted, the Cloud Speech API can perform **speaker diarization**, labeling transcript segments by speaker in multi-speaker audio. This is achieved by an integrated speaker embedding model that runs alongside transcription. The Gemini system, being conversational, is *context-aware* enough to distinguish the user's voice from its own responses, and it is trained to ignore background voices or interruptions ("conversation context awareness") – effectively a form of speaker/activity detection to decide when to listen and when to talk. DeepSeek, if used in a voice assistant capacity, would similarly need to integrate voice activity detection and possibly speaker recognition to function in the real world, though specifics are not published. Audio classification (like detecting non-speech sounds or classifying music genres) is not a primary feature of Whisper or the language-centric models, but large multimodal models can be prompted to categorize sounds if given an audio input. For instance, one could transcribe audio with Whisper or USM and then ask an LLM (like GPT-4 or DeepSeek) to interpret the content (e.g., "the sound of applause" or "classical piano music"). Google Gemini's multimodal training likely included some capacity to interpret audio events – as evidenced by claims that it can handle "*subtle auditory cues*" in a conversation – though how far this extends beyond speech is unclear. In summary, for audio understanding tasks, specialized models (Whisper, USM) currently lead in transcription performance, while larger AI systems incorporate these capabilities to form part of a bigger reasoning process. Table 1 (in the Comparative Analysis section) will summarize some features of these models.

Technical Approaches to Real-Time Processing

An important aspect of audio understanding is whether it can be done **in real time**. Real-time ASR requires processing streaming audio with minimal delay (typically a few hundred milliseconds). This often necessitates architectural tweaks: for example, using *online* models like Recurrent Neural Network Transducers (RNN-T) or chunk-based processing with Transformers. Whisper, as originally released, processes audio in fixed 30-second blocks, introducing a delay equal to the block size. However, developers have adapted Whisper for streaming by using shorter chunks (e.g., 5-second windows) at slight cost to accuracy. OpenAI's newer Whisper API model ("Whisper-1") and especially the experimental **Whisper large-v3 Turbo** show improved speeds – one community report indicates the Whisper large-v3-turbo model can transcribe nearly in real-time with ~0.3 s latency on appropriate hardware. Google's ASR models have been designed for streaming from the outset: the RNN-T based models that powered Google Assistant could output words as you speak, and the latest Conformer-based **Chirp 2** model supports streaming recognition

for real-time applications ² . These models use lookahead and chunking to balance accuracy and latency. DeepSeek’s full model, due to its sheer size, is not geared for low-latency processing of audio; instead, one would use a faster ASR front-end (possibly a smaller DeepSeek “Turbo” model or an external API) to handle real-time transcription, then feed results into DeepSeek’s main model for higher-level analysis or response. In Section **Real-Time and Offline Capabilities**, we further compare these systems’ ability to operate under real-time constraints.

Technical Background: AI Models for Audio Generation

Text-to-Speech (TTS) and Voice Cloning Architectures

On the generation side of audio, **text-to-speech** (TTS) models convert written text into spoken voice. Traditional TTS systems involved a pipeline of text front-end (linguistic analysis, phonetic transcription, prosody prediction) and a back-end vocoder to produce audio waveforms. Modern AI-based TTS has shifted to end-to-end neural architectures that can learn to synthesize speech directly from text, given enough training data of human recordings. A landmark was **WaveNet** (DeepMind, 2016), a generative model for raw audio that achieved unprecedented naturalness in synthetic speech, with Mean Opinion Scores (MOS) approaching 4.2 out of 5 (where human speech is ~4.7). WaveNet, however, was slow due to sample-level generation. Subsequent models (e.g., Tacotron 2 + WaveNet, FastSpeech, etc.) improved efficiency by generating intermediate acoustic features (like mel-spectrograms) which are then converted to audio by neural vocoders.

OpenAI’s **Voice Engine** represents the cutting edge of TTS combined with *voice cloning*. According to OpenAI, Voice Engine is a “small model” that can produce natural, emotive speech in a voice that closely resembles a given example recording. In a preview announced in 2024, OpenAI revealed that Voice Engine allows a user to upload a **15-second sample** of a target speaker’s voice and then generates speech in that voice from any text prompt. This capability is similar to research like Microsoft’s VALL-E, which demonstrated zero-shot voice cloning from a few seconds of audio. The model underlying Voice Engine had been in development for two years and was hiding “in plain sight” – it actually powers the preset voices of OpenAI’s text-to-speech API and the voice output in ChatGPT’s mobile app. In terms of architecture, OpenAI hasn’t published detailed specifics, but we can infer it uses a sophisticated **neural codec** or **embedding** approach to capture voice characteristics. It likely encodes the 15s sample into a speaker embedding, and uses a multi-speaker TTS model (possibly an auto-regressive or diffusion-based model operating on audio codec tokens) to generate new speech in that voice. The fact that it’s described as a “small model” suggests OpenAI found an architecture efficient enough to run relatively quickly (important for real-time applications like ChatGPT’s responses). The speech quality is reportedly highly realistic – the cloned voices preserve accent, tone, and prosody of the original speaker, even expressing emotions to some degree. Critically, OpenAI has **not** widely released Voice Engine to the public as a general tool, due to concerns about misuse (e.g., impersonation and deepfakes). Instead, they have run a limited preview with trusted partners and implemented safeguards: requiring explicit consent from any individual whose voice is cloned, prohibiting certain use cases, and embedding an imperceptible watermark in generated audio to trace its origin. These safety measures reflect the ethical challenges of voice cloning, which we discuss later. From a technical standpoint, Voice Engine demonstrates that *few-shot voice cloning* is now feasible with high fidelity – an achievement made possible by combining large training datasets of diverse voices with advanced generative modeling. OpenAI’s work in this area is a significant contribution, as it promises more natural and personalized TTS systems (e.g., allowing users to have an AI speak in their own voice for assistive technology).

Google's TTS capabilities are likewise very advanced and have been integrated into the Gemini ecosystem. Google has a history of research in expressive TTS, from WaveNet to Tacotron to more recent models like **AudioLM** and **VoiceFilter**. In Gemini 2.5, Google introduced *controllable* TTS as a native feature. This means developers can prompt Gemini to generate speech with specific styles, tones, or emotions using natural language instructions. The system can produce anything from a neutral newscaster voice to a lively storytelling tone, or mimic accents and even whisper, all at the user's request. Gemini's TTS supports **multilingual** output (24+ languages) and can seamlessly switch languages or do code-switching in one utterance. A striking feature is *multi-speaker dialogue generation*: Gemini can take text for a conversation between two characters and generate an audio clip where two distinct voices perform the dialogue, making the output more engaging. The underlying technology likely builds on Google's research in cascaded audio models – possibly using a Transformer-based language model to generate semantic or acoustic tokens which are then rendered to speech via a learned codec (Google's AudioLM, for instance, uses SoundStream codec tokens and Transformers to generate coherent speech with long-term consistency). Google has reported that their models move “beyond naturalness” to give unprecedented **control** over style and delivery. Unlike OpenAI's Voice Engine, Google's public offerings do not currently allow arbitrary voice cloning from user-provided samples – instead, Google provides many preset high-quality voices and the ability to modulate them. This difference stems from policy rather than technical inability; Google no doubt has the technical means for voice cloning (they demonstrated Speaker Adaptation in TTS research), but like OpenAI, they are cognizant of misuse risks. One important contribution from Google is the development of **SynthID** for audio, a watermarking technology embedded in all AI-generated audio from Gemini 2.5. This allows detection of AI-generated speech to help combat deception. In summary, Google's Gemini approach to speech generation emphasizes integration (TTS as part of a larger assistant that can think and then speak) and fine-grained controllability in real time. Latency has been optimized to be “*very low*” for fluid conversation – likely by using a faster decoding model (Gemini 2.5 has a “Flash” mode for speedy responses, in addition to a “Pro” mode for highest quality).

DeepSeek has been slightly less public about any text-to-speech or voice cloning capabilities. The DeepSeek models are primarily known as text generators and multimodal analysts. However, given DeepSeek's goal to be a general AI, it could be extended to generate audio as well. There is no indication that DeepSeek can clone arbitrary voices at present. It may rely on external TTS engines to speak its responses (for example, a user could use DeepSeek to produce a text answer and then use a third-party TTS to vocalize it). That said, one of DeepSeek's headline features is **music generation**, which is a form of audio generation. DeepSeek-V3 has been highlighted for its ability to generate music compositions with AI. This suggests DeepSeek trained a specialized component (or expert) on music data, potentially similar to how OpenAI's **Jukebox** model works or Google's **MusicLM**. Jukebox (OpenAI, 2020) was a pioneering model that generated songs with lyrics by modeling audio at the waveform level with a hierarchical VQ-VAE and Transformers. Jukebox was extremely computationally heavy – taking about *3 hours to generate 20 seconds* of music on a V100 GPU – and its output, while often coherent in style, could be hit-or-miss in musical quality. By 2023, Google introduced **MusicLM**, a text-to-music model that produces 24 kHz audio consistent with a given description (and can even be conditioned on humming) ³. MusicLM outperformed previous systems in quality and could maintain musical structure over several minutes ⁴. However, Google initially did not release MusicLM publicly due to concerns about copyright (the model was trained on thousands of hours of music, potentially memorizing riffs). In 2024, DeepSeek's team seemingly managed to create an AI music model at least on par with those Western efforts. The DeepSeek-V3 music generation model (and its open-source distilled versions like R1-Distill) can be deployed locally, offering both cloud-based and offline generation of music. While details are scant, it likely uses a combination of *massive model size* and *mixture-of-experts* to capture different facets of music (melody, harmony, instrumentation). The significance is that DeepSeek's

approach leverages its enormous parameter count to handle audio generation tasks that typically require specialized models. Performance-wise, the music generated by DeepSeek has been described as a “symphony of disruption” to the industry, indicating high quality and prompting discussions on how it might affect musicians and copyright.

In summary, AI audio generation has reached a point where **speech synthesis is nearly indistinguishable from human speech** and can be personalized (OpenAI and Google’s contributions), and **music synthesis is becoming realistic** (Google’s and DeepSeek’s contributions). The architectures enabling this include large Transformer models operating on learned audio representations, often with the help of discrete *audio codes* (from neural codecs) to manage the high data rate of audio. These models can capture not only the content of speech or music but also style and timbre, allowing for emotive and context-appropriate generation. They increasingly operate in **real-time or near-real-time**, making interactive voice AI feasible. The next sections will compare these platforms on key metrics and discuss how they perform in live settings (real-time vs offline) and various use cases.

Comparative Analysis of OpenAI, Google Gemini, and DeepSeek Approaches

To elucidate the strengths and trade-offs of OpenAI’s, Google’s, and DeepSeek’s audio AI systems, we compare them along several dimensions: **model performance, latency, data efficiency, and scalability & deployment**. Table 1 provides a high-level comparison of key features for each platform, and we discuss each aspect in detail below.

Table 1. Key Features of AI Audio Platforms – OpenAI vs Google (Gemini) vs DeepSeek

Aspect	OpenAI (Whisper & Voice Engine)	Google Gemini (incl. USM & TTS)	DeepSeek (R1/V3 models)
Speech Recognition (ASR)	Whisper transformer (open-source); ~1.5B params (large); 680k hours supervised data. Multilingual (98 languages), robust to noise; ~5% WER on English test, not SOTA on LibriSpeech. Real-time with Whisper-turbo (~0.3s latency) on GPU.	USM Conformer encoder (2B params) with RNN-T/CTC decoder. 12M hours self-supervised + labeled fine-tune; 100+ languages. SOTA accuracy: ~30% avg WER on 73 languages, outperforming Whisper by ~33% on overlap. Supports streaming ASR with low latency (production used in Assistant/YouTube).	Likely uses external or integrated ASR front-end (DeepSeek can transcribe via API). Multimodal model (37B active params) can process audio input, but heavy for direct ASR. Distilled models (R1) available for local use. Accuracy not benchmarked publicly; depends on underlying ASR component (possibly Whisper or similar).

Aspect	OpenAI (Whisper & Voice Engine)	Google Gemini (incl. USM & TTS)	DeepSeek (R1/V3 models)
Speech Generation (TTS)	Voice Engine TTS (model size not disclosed, described as "small"). Clones voice from 15s sample; highly natural and emotive output closely resembling speaker. Not publicly available yet (limited preview); used in ChatGPT voice (5 voices) and API. Near real-time generation; outputs watermarked audio for safety.	Gemini 2.5 native TTS with controllable style. High-quality neural voices (WaveNet-level or better); supports 24+ languages. Can adjust prosody, emotion, accent via prompts. Can generate multi-speaker conversations in one go. Very low latency speech generation for interactive use; all AI audio outputs embedded with SynthID watermark.	No standalone TTS product known. DeepSeek not focused on TTS for user output; likely relies on third-party TTS to vocalize responses. The large model could in principle generate speech via an audio token decoder (not confirmed). No voice cloning feature announced.
Music/Sound Generation	OpenAI Jukebox (2020) for music; 5B param model, produced 1+ min songs but very slow (hours of compute for seconds of audio). Not real-time or practical for deployment. No recent public music model (OpenAI focuses on speech over music).	MusicLM (research prototype) generates music from text, 24 kHz, multi-minute coherence ³ . Not integrated into products yet (experimental in AI Test Kitchen). Gemini not explicitly a music model, but multimodal capabilities could extend to audio events. Google has separate models for sound effects, but Gemini mainly targets speech audio dialog.	DeepSeek V3 includes AI music generation as a core feature ¹ . Models (e.g., DeepSeek-R1-Distill) can produce music locally. Quality reported as industry-leading, causing discussions on copyright and creative impact. Not clear if real-time (likely not, given complexity). No known model specifically for sound FX generation, but the model could be prompted to create simple sounds (requires an audio output mechanism).

Aspect	OpenAI (Whisper & Voice Engine)	Google Gemini (incl. USM & TTS)	DeepSeek (R1/V3 models)
Latency (Real-Time)	Whisper: initially offline batch (30s segments) – ~2x real-time on CPU for large model, faster on GPU. New Whisper-3 Turbo can do ~real-time streaming. Voice Engine: designed for interactive use (ChatGPT responds in seconds), so likely <1s synthesis for a sentence (with GPU inference). Overall OpenAI: real-time feasible with optimized models/hardware, but Whisper still slower than some optimized ASR.	Google: optimized for streaming. ASR (USM/ Chirp) can transcribe with <300ms delay in cloud. Gemini’s dialog system runs in real-time: processes input and speaks reply with very low latency. Google’s TTS can generate speech faster-than-real-time (Cloud TTS can synthesize 1s audio in <1s). Thus, Gemini achieves smooth turn-taking in conversation.	DeepSeek: full 671B MoE model is too large for real-time on normal hardware. Inference requires multi-GPU or TPU, causing significant delays for long inputs. Smaller DeepSeek models (R1 Turbo) improve speed but still not as fast as dedicated ASR/ TTS. Likely scenario: not real-time for direct transcription; uses streaming ASR front-end and only the LLM for understanding responses (incurring a few seconds for complex reasoning). Not suitable for low-latency voice dialog without substantial optimization.

Aspect	OpenAI (Whisper & Voice Engine)	Google Gemini (incl. USM & TTS)	DeepSeek (R1/V3 models)
Data Efficiency & Training	Relies on huge labeled data for ASR (680k hr). Whisper was not trained with self-supervision, so it required careful dataset curation. For TTS, OpenAI used a mix of licensed/public voice data; details scarce due to IP concerns. Voice Engine's success with only 15s sample indicates efficient few-shot learning of voice – likely due to rich training covering many speakers. OpenAI tends to favor large-scale training for generalization (Whisper wasn't fine-tuned to specific domains, yielding broad robustness).	Massive unsupervised pre-training (12M hr) allowed state-of-the-art results with less labeled data for new languages. Google's models thus achieve better low-resource language performance (self-supervised learning for scale). For TTS, Google benefits from large internal datasets of human recordings in many languages (e.g., WaveNet was trained on thousands of hours). Gemini being multimodal means it draws on varied data (text, audio, images); this diversity helps it generalize but also requires colossal data and compute. Google's approach is data-hungry but yields strong general models that can be adapted (e.g., fine-tuning a telephony ASR on phone audio).	DeepSeek's strategy is to throw enormous model capacity at the problem and integrate diverse data. DeepSeek-V3's 671B MoE was trained on vast data across modalities (the exact mix isn't public, but likely hundreds of billions of tokens of text and a large corpus of images/audio). Data efficiency is achieved via MoE – experts specialize, so not every parameter sees all data. This, plus continual learning (DeepSeek updates models continuously), means DeepSeek can ingest new information without full retraining. However, the initial training would have been extremely data and compute-intensive. In terms of few-shot abilities, DeepSeek can interpret new audio tasks on the fly if within its learned capability, but heavy fine-tuning is not feasible on the giant model (hence the need for distilled smaller models for practical use).

Aspect	OpenAI (Whisper & Voice Engine)	Google Gemini (incl. USM & TTS)	DeepSeek (R1/V3 models)
Scalability & Deployment	OpenAI provides cloud APIs for Whisper (making heavy models accessible without users running them). Whisper is also open-source, so it's deployed on devices and servers by many – the small models run on laptops or smartphones (with quantization). Voice Engine is currently only cloud-based (and only for select partners), not deployable by end-users due to misuse risk. Scalability-wise, OpenAI's approach favors <i>specialized models</i> for each task (ASR, TTS) that can be integrated – this modular design can be scaled by improving each component. OpenAI has also leveraged model compression for deployment, e.g., Whisper tiny for edge devices.	Google scales by offering models at multiple sizes (e.g., “Base, Large, or Chirp” models for Cloud ASR) and by high-performance inference on Google Cloud TPUs. Gemini itself has <i>tiers</i> (Gemini 2.5 Pro vs Flash) to balance quality vs speed. For deployment, Google has started bringing some advanced models on-device: e.g., some Pixel phones run a version of speech model for voice typing offline. Generally, though, full Gemini runs in cloud due to its size. Google's strength in scaling is the integration – one Gemini model (or family) handles many tasks, simplifying deployment for multiple capabilities (but also concentrating risk if it fails). Scalability in training is achieved with distributed infrastructure (TPU pods) that only a few organizations possess.	DeepSeek's largest models are not easily deployable except on powerful clusters, which the company offers via API or private instances. However, DeepSeek made smaller variants (DeepSeek Lite, R1-distill) for <i>local</i> use, acknowledging the need for accessibility. These can run on a single high-end GPU, enabling offline use by enthusiasts or enterprises with hardware. The MoE design is inherently scalable – more experts can be added to improve performance without increasing per-token computation drastically. This means DeepSeek could further scale to even larger parameter counts. From a deployment view, DeepSeek positions itself as an all-in-one model (answering questions, generating content, analyzing data), which users access via a chat interface or API. This monolithic approach is powerful but can be inefficient if only a narrow function (like just ASR) is needed. Thus, for pure audio tasks, DeepSeek may be overkill, but for multimodal tasks that include audio, it shines by not requiring separate systems.

Performance and Accuracy: In terms of raw performance on core tasks, Google's specialized models (as incorporated in Gemini) currently lead in **ASR accuracy**, especially for a wide range of languages. As noted, Google's USM achieved a significant WER reduction over Whisper on many multilingual benchmarks. Whisper remains excellent (often near human-level on clear English speech) and has the advantage of being openly available, but Google's investment in unlabeled data and larger training sets yielded a more universally accurate model by 2023. For **speech generation quality**, both OpenAI's Voice Engine and

Google's TTS produce remarkably human-like speech. It is hard to declare a "winner" as both can achieve MOS scores very close to human recordings (likely in the 4.5+ out of 5 range in internal tests). OpenAI's voice cloning is a unique strength – generating a custom voice that mimics a specific person is something Google hasn't offered publicly (Google focuses on expressive preset voices). This means OpenAI's system can personalize more, which is a performance aspect beyond numeric scores – *speaker similarity* is a metric where Voice Engine excels (cloned speech "closely resembles the original speaker"). For **music generation**, OpenAI's early work (Jukebox) showed what was possible, but it's DeepSeek that currently claims the spotlight with its V3 music model. Reports suggest DeepSeek-V3 can generate coherent music pieces that impress listeners, potentially matching or surpassing Google's unreleased MusicLM in quality. However, without standardized benchmarks (beyond subjective listening tests) it's difficult to quantify. It's clear all three players have pushed audio generation performance: music from text prompts (DeepSeek, Google) and voice from text with style (OpenAI, Google) are now at a level that was science fiction a decade ago.

Latency and Real-Time Capability: Google's Gemini is explicitly engineered for **real-time dialog** – this is a major selling point of the model. Google researchers highlight "*very low latency*" speech interactions and fluid conversation with Gemini 2.5. This is achieved by efficient streaming ASR (likely processing audio in chunks of a few hundred milliseconds) and fast speech synthesis (the "Flash" TTS model trades off some fidelity for speed to ensure quick responses). Therefore, for use cases like voice assistants or live translations, Google's platform currently has the edge. OpenAI's Whisper, out-of-the-box, introduces noticeable latency due to its chunk size (e.g., waiting up to 30s before output). But with modifications (smaller chunks, or the Whisper API model which might use a sliding window technique), OpenAI has brought latency down to a more interactive regime. For instance, if Whisper can achieve ~0.3 seconds latency for a short utterance on good hardware, that is effectively real-time for many purposes. OpenAI's Voice Engine in ChatGPT responds in near real-time (the system begins speaking output while still formulating later parts of the text). This suggests OpenAI uses streaming output: possibly GPT-4 generates text which is then synthesized on the fly by the Voice Engine. The voice synthesis portion is fast enough that you hear only a minor pause before the AI speaks. In short, OpenAI *can* support real-time experiences, but it required optimizing their models and infrastructure. DeepSeek, by virtue of its size, is the least real-time-friendly. DeepSeek might be more suited to offline or *asynchronous* processing – for example, transcribe a meeting recording and summarize it after the fact, rather than transcribing live with minimal delay. Even so, if one uses DeepSeek-R1 Turbo (a smaller model) and a separate streaming ASR, one could build a pseudo-real-time pipeline: ASR produces text live, DeepSeek takes the text and quickly generates a response or action. But compared to OpenAI and Google, DeepSeek has not demonstrated an equally snappy interactive voice loop. It's more akin to a powerful back-end brain that could be used in a voice assistant if coupled with faster front-end components.

Data Efficiency: This measures how well each approach makes use of data, especially low-resource scenarios. Google's approach shines here: by using self-supervised pre-training, Google achieved strong results for languages with very little transcribed data. This is a clear efficiency win – rather than needing hundreds of hours of labeled audio for each new language, USM could learn from raw audio. OpenAI's Whisper took a brute-force supervised approach; it works fantastically for the languages it was trained on (which fortunately include many major languages), but if you ask Whisper to transcribe a language it hasn't seen much, it might fail. OpenAI has not (yet) implemented techniques like iterative pseudo-labeling or self-supervision for Whisper, so in that sense Google's approach is more data-efficient and scalable to new domains. For TTS, OpenAI's voice cloning is extremely efficient in terms of speaker data – 15 seconds of audio is enough to capture a voice's essence. Achieving this implies the model generalizes well from prior training on possibly thousands of voices. Google's TTS typically would require on the order of minutes of

audio (or adaptation training) to create a new custom voice, since they haven't offered a 15-second cloning feature. So OpenAI demonstrated greater efficiency in *speaker adaptation*. DeepSeek's data efficiency is harder to judge – it uses enormous data and parameter count, which is the opposite of frugality. However, by covering so many tasks in one model, one could argue DeepSeek is data-efficient at *multitask learning*: it doesn't need separate models (and separate datasets) for ASR, translation, summarization, etc., but rather one ginormous training process. For an individual task, though, it likely uses more data and compute than the others to reach a given level of performance.

Scalability and Deployment: OpenAI's strategy has been to open-source core models like Whisper to spur widespread adoption (and indeed Whisper is now integrated in countless applications from transcription services to homebrew voice assistants). This openness means OpenAI's tech can scale **horizontally** through the community. At the same time, OpenAI provides **hosted APIs** (for Whisper and for TTS) which let developers use these models without running them locally. This cloud API approach is scalable in that OpenAI can upgrade the model behind the scenes (like releasing Whisper "large-v2" or a faster variant) and all users benefit. A limitation for OpenAI is that Voice Engine is not generally released – scaling its impact is slowed by the cautious rollout (only select partners have it). Google, by contrast, often integrates new models directly into consumer-facing products (Android's GBoard for speech-to-text, YouTube for captions, etc.), instantly reaching millions of users. With Gemini, Google is unifying capabilities, which can simplify deployment: instead of maintaining separate systems for ASR, assistant dialogue, and TTS, Google can use one model (or a coordinated set of models under the Gemini umbrella). However, running a multimodal behemoth like Gemini is costly; Google likely will use smaller variants or distillations for lightweight applications (e.g., on-device or free-tier services), reserving full Gemini for premium services. DeepSeek, being a newer player, has smartly recognized the need for accessible deployment: by releasing **DeepSeek R1 Distill** open-source, they allowed the community to experiment locally (albeit with a downsized model). DeepSeek also offers a web platform and API (with both free and enterprise options), which is how they presumably manage the heavy lifting for users who need the full power. The MoE design of DeepSeek aids scalability – one can increase expert count to improve performance without linearly increasing inference cost. This means DeepSeek could scale up to even more parameters if needed (given enough hardware), and theoretically handle more simultaneous tasks by routing to different experts. In practical terms, OpenAI and Google have the advantage of massive infrastructure and user base to deploy their models widely, whereas DeepSeek is still building out its ecosystem. However, DeepSeek's rapid progress (from founding in 2023 to having multiple large models by 2025) suggests it is poised to compete by offering its technology perhaps at lower cost or with open-source transparency, which can drive adoption.

In conclusion, each platform has distinct advantages: **OpenAI** offers *specialized, high-quality models* (like Whisper and Voice Engine) with an emphasis on open research and cautious deployment; **Google (Gemini)** leverages *integrated multimodal AI* and a track record of production-grade speech services to excel in real-time, broad-coverage scenarios; **DeepSeek** pushes the envelope on *scale* and *multimodal learning*, potentially leapfrogging in areas like music generation and unified AI tasks, though with challenges in efficiency. The comparative insights gleaned here will inform our discussion of specific use cases and where each platform might be best applied.

Real-Time and Offline Capabilities

A critical practical consideration for audio AI is whether it can run **offline** (without internet/cloud) and whether it can operate in **real-time** for live interactions. We examine each platform with respect to these capabilities:

- **OpenAI (Whisper & Voice Engine):** Whisper’s open-source release means it can be run offline on local hardware. Many users run Whisper on laptops or even smartphones (the smallest Whisper models) to transcribe audio privately. The trade-off is speed: Whisper-large, with ~1.5B parameters, typically requires a GPU to run faster than real-time. On CPU, real-time transcription is achievable only with the tiny or base models, or by tolerating some delay. Nonetheless, the ability to run completely offline is a major advantage for privacy-sensitive applications (e.g. transcribing meetings or voice notes without sending data to the cloud). OpenAI’s Voice Engine, on the other hand, is not available for local use – it’s only on OpenAI’s servers for now. This means voice cloning with OpenAI tech currently requires cloud access and is limited to approved partners. For real-time use, as discussed, OpenAI has optimized its pipeline so that ChatGPT with voice feels interactive. When you speak to the ChatGPT mobile app, your speech is transcribed (likely by Whisper or a variant) nearly instantly, and GPT-4 responds; the response text is then spoken by Voice Engine with minimal delay. This whole loop is near real-time but does involve cloud processing at each step (the app streams your voice to the server). If connectivity is lost, these features would not function since the heavy models aren’t running on the device. In summary, OpenAI enables **offline use for ASR** (via Whisper open source) but **not for voice cloning/TTS** (Voice Engine is cloud-only). Real-time is achievable with OpenAI’s cloud (they demonstrate it in ChatGPT voice conversations), but using Whisper offline in real-time may require powerful local hardware or reduced model sizes.
- **Google Gemini (and Cloud Speech services):** Google’s ecosystem historically offers both cloud and some offline options. For example, GBoard’s voice typing can run on-device for a few languages using a small recurrent model. The new Gemini audio capabilities, however, are currently cloud-centric – they are accessed via Google’s AI Studio or API endpoints. Gemini is far too large to run on common personal devices in its full form. Instead, Google might deploy smaller distilled models to devices in the future (just as they did with earlier voice models). At present, we can consider Gemini’s advanced audio understanding and generation as an **online service**. This means an internet connection to Google’s servers is needed to use features like real-time multilingual conversation or controllable TTS. Google’s Cloud Speech-to-Text API does have an **on-premises** version and even an offline SDK (for embedded usage, Google released “Speech On-Device API” for some enterprise applications), but those use older models (not the latest Gemini/USM). The *real-time* performance of Google’s cloud is excellent: many developers use Google’s streaming ASR API to transcribe audio on-the-fly with only a few hundred milliseconds of latency. Likewise, Cloud TTS can stream audio as it’s generated. For instance, when using Google Assistant, after you finish speaking a question, the assistant typically responds almost immediately; behind the scenes, Google’s ASR may only need ~200 ms after you stop talking to finalize the transcript (due to endpointing algorithms), and the response is generated and spoken with perhaps ~500 ms of additional delay. These sub-second timescales make the interaction feel fluid. Gemini 2.5’s own demonstrations at Google I/O 2025 showed live conversation where the system would speak with “*remarkable quality... delivered with very low latency*”. Therefore, for real-time applications (like interpretation, live assistance, etc.), Google’s platform is very well-suited *when online*. If offline operation is required, Google’s most advanced models are less accessible than Whisper – one would likely fall back to smaller models (which have

lower accuracy or less functionality). In summary, **Gemini is a real-time pro** in connected scenarios, but **offline** it offers only limited, less powerful alternatives at present.

- **DeepSeek:** Given its resource demands, DeepSeek's full models are predominantly cloud-hosted. DeepSeek AI is accessible via a web interface and APIs, meaning your data goes to their servers for processing. However, DeepSeek did release lighter versions like "DeepSeek Lite" and provided means to run models like R1-Distill locally. Running these offline still requires a decent GPU and a lot of memory, but it's feasible for some enthusiasts or organizations. DeepSeek's openness in this regard is noteworthy – while the largest models remain proprietary, they encourage community adoption through slightly smaller open models. In offline mode, one obviously cannot leverage the 671B-param capacity, but perhaps a 6B or 13B param variant could be used, which might be enough for many tasks but with reduced accuracy. **Real-time** use of DeepSeek is where it struggles the most. A user query to the 37B MoE model might take a few seconds to process even on a good server (and longer if the query is very complex, since these LLMs often have high per-token generation times). DeepSeek is likely too slow to perform instantaneous transcription or to generate immediate voice responses without perceptible lag. If building a voice assistant with DeepSeek, one strategy would be to use a fast ASR (like Whisper small or Google's API) to get text, then use DeepSeek for the **reasoning** or understanding part, then a fast TTS to respond. DeepSeek would effectively function like a powerful but somewhat slower brain in the loop. Users might tolerate a few seconds of delay for a complex response (similar to how some use GPT-4 via API with a short wait). But for rapid back-and-forth dialogue, DeepSeek in 2025 is not ideal. Its strength lies in processing tasks that don't require instantaneous turnaround – e.g., analyzing an hour-long audio recording and producing a summary is something that can take a minute or two on DeepSeek, which is fine. But taking turns in a live conversation with <1s latency is beyond DeepSeek's current deployment. Perhaps in the future, optimizations or hardware advances (or the use of MoE to route tasks to smaller experts) could enable faster response. For now, we consider **DeepSeek a cloud-oriented solution with slower real-time performance**, and partial offline usability via smaller models.

In conclusion, if an application demands fully **offline** functionality (no internet), OpenAI's Whisper is the most practical for recognition, and one might pair it with an offline TTS engine from other sources (since OpenAI's best TTS isn't available offline). Google's older on-device models can handle simple voice commands offline, but for anything like the full Gemini experience, internet is required. DeepSeek offline is possible only with significantly reduced models that may not outperform the alternatives. If **real-time interactivity** is the goal (say, a conversational agent or real-time translation device), Google's platform currently offers the best end-to-end latency and smoothness, with OpenAI a close second when using its cloud (as evidenced by the new voice mode in ChatGPT). DeepSeek would rank third, suitable for near-real-time at best (short delays), but not instantaneous responses. The choice often boils down to the use case: for a live bilingual translator, one might lean on Google; for a privacy-sensitive transcription tool with no internet, one would use Whisper; for an in-depth analysis of audio content (like extracting insights from a recorded meeting), one might accept a slower DeepSeek pipeline for its advanced understanding capabilities.

Applications of AI Audio Technologies

AI models for audio understanding and generation have already begun transforming numerous industries and user experiences. We highlight some key real-world applications and how OpenAI's, Google's, and DeepSeek's technologies contribute:

- **Voice Assistants and Conversational AI:** Perhaps the most visible application is in virtual assistants (like Siri, Alexa, Google Assistant, and now ChatGPT's voice mode). These systems rely on robust ASR to understand user queries, natural language understanding to interpret and reason, and TTS to reply in a friendly voice. Advances by OpenAI and Google directly enhance these assistants. For example, Google Gemini's real-time dialog capabilities will likely power the next generation of Google Assistant, allowing it to handle complex queries with reasoning and respond fluidly in multiple languages. OpenAI's technology enables **ChatGPT with voice**, where Whisper transcribes what the user says and GPT-4 plus Voice Engine respond with a human-like voice. This opens up interactive chat with an AI as a hands-free, eyes-free experience – users can have a conversation with an AI as if they were talking to another person. DeepSeek could be used in specialized assistants that need higher-level cognitive functions – for instance, an AI doctor's assistant that listens to a patient consultation (transcribed by Whisper or USM) and then uses DeepSeek to suggest diagnoses or questions. In customer service, these models can drive voice bots that handle calls, providing more natural interactions. Speaker identification can also be used in this context for personalization (recognizing a user's voice to retrieve their preferences or authenticate identity, although caution is needed as voice can be spoofed in the era of cloning).
- **Transcription and Media Captioning:** A straightforward but impactful application of ASR is transcribing audio and video content into text. This is useful for creating subtitles for videos, generating meeting minutes from discussions, or enabling search within audio archives (like podcasts or call logs). OpenAI Whisper's high accuracy and open availability have made it a go-to tool for many media professionals – for instance, journalists transcribing interviews or YouTubers auto-generating captions for their videos. Google's ASR (through services like YouTube's caption auto-generation and Google Meet live captions) has scaled this to millions of users, providing real-time captions that aid accessibility (e.g., for the hearing-impaired or multilingual audiences). With models like USM, these captions are more accurate and available in more languages than ever. DeepSeek's role here could be to *enhance* transcription outputs – for example, DeepSeek could take a raw transcript of a meeting and produce an abstractive summary or identify key action items (combining its language understanding with the transcription). Indeed, workflows have been proposed where Whisper does the transcription and DeepSeek (R1 LLM) summarizes the result. This division of labor plays to each system's strength. Overall, transcription automation saves countless hours of manual typing and enables content to be consumed in different forms (audio turned to text).
- **Content Creation and Dubbing:** AI audio generation is enabling new forms of content creation. One example is **dubbing podcasts or videos into other languages** using the same speaker's voice. OpenAI's Voice Engine, in partnership with companies like Spotify, has been used to translate popular podcasts: a host speaks in English, and the AI produces the same speech in Spanish or other languages, preserving the host's own voice timbre. This is revolutionary for media localization – audiences can hear content in their native language without losing the personal voice of the creator. Similarly, startups like HeyGen have used Voice Engine to generate multilingual video voice-overs for

corporate content. These uses rely on both the transcription (to get the original script) and the voice cloning TTS (to speak it in another language). **Audiobook narration** is another emerging application – an AI could clone a voice (even a famous actor’s voice, with permission) to narrate books, saving time in recording studios. On the visual side, generating **custom AI avatars** with synchronized lip-synced speech (text-to-video with audio) is becoming possible; the audio part is driven by these TTS models (Gemini, for instance, could be used to give a voice to a generated avatar speaking any language on the fly). DeepSeek’s ability to generate long-form content (like extended discussions or music backgrounds) can support content creators in prototyping or enhancing creative works. For instance, a game developer could use DeepSeek to generate ambient soundscapes or theme music based on a textual description of a scene, potentially speeding up the creative process.

- **Music and Entertainment:** AI’s foray into music generation creates new opportunities in the music industry and beyond. A musician can use models like MusicLM or DeepSeek to **generate musical ideas or accompaniments**. For example, someone with a melody in mind but lacking an orchestra could hum the tune and have the AI produce a rich orchestral interpretation in a given style ⁵. DeepSeek’s model, which has been highlighted in media as a “wake-up call” for the music industry, might allow rapid creation of music for commercials, video games, or films without needing a full composition team – though not without raising copyright and creative authenticity questions. On the consumer side, these models might enable personalized music (an AI DJ that generates songs on the fly tailored to a listener’s mood). We also see AI being used to **remix or restore audio** – e.g., separating music tracks, changing the instrumentation or style of a song (transfer learning on audio). While not explicitly covered above, these are tasks that can be aided by AI models that understand audio content. Google’s AudioLM was shown to continue piano pieces in the style of the input, demonstrating how AI can extend or fill in music, a tool for musicians during editing.
- **Healthcare and Accessibility:** Audio AI has meaningful applications in healthcare. Speech recognition can power medical dictation software (doctors speaking clinical notes that get transcribed automatically, saving them paperwork time). More advanced is analyzing patient speech for signs of conditions – for example, changes in voice can signal Parkinson’s disease or stroke; cough sounds can help diagnose respiratory illness. An AI model trained on health data could flag these from audio. On the generation side, Voice Engine and similar tech can **restore voices** for people who have lost the ability to speak (one of OpenAI’s pilot projects did this for a patient, using a prior recording of her voice ⁶). This is life-changing for people with conditions like ALS – they can communicate in their own reconstructed voice via a text-to-speech system, rather than a generic robotic voice. Similarly, for the visually impaired, AI can serve as a reading assistant: models like Voice Engine can narrate text with human-like intonation, and even take on a voice that the user finds comforting or familiar. There are also applications in mental health – an AI that listens (with consent) to therapy sessions could help analyze a patient’s emotional state from vocal cues, or provide transcription for review. While these are emerging and need careful validation, they illustrate the broad societal impact of audio AI beyond tech gadgets.
- **Security and Forensics:** Audio AI can be a double-edged sword in security. On one hand, speaker recognition can be used for biometric authentication (voice ID to unlock devices, verify callers in banking, etc.). However, given voice cloning advances, using voice as a password has become riskier (attackers could synthesize the target’s voice). Indeed, OpenAI explicitly recommends phasing out voice-based authentication because of how convincing generated voices have become. Instead, AI can help *detect* fraud – for example, voice anti-spoofing models can try to tell if a voice is real or AI-

generated (by catching artifacts that humans can't hear easily). Companies are working on detector models in tandem with generation models (e.g., the mentioned SynthID watermark from Google, or audio watermarking by OpenAI). In law enforcement forensics, AI transcription is invaluable for analyzing wiretaps or recordings, and diarization helps know who said what. DeepSeek's large-scale analysis could potentially cross-reference audio content with text knowledge bases (for instance, identifying if a certain quote in an audio clip matches something from a database). Privacy tools might also use AI to **anonymize voices** – by changing the speaker's voice in an audio (while preserving words) to protect identities, which is essentially a generation task (voice conversion). AI can facilitate such transformations in high quality, which could help in journalistic interviews or witness protection scenarios.

These applications illustrate the versatility of the audio AI technologies developed by OpenAI, Google, and DeepSeek. Many of these uses are already in deployment (automatic captions, virtual assistants), while others are just on the horizon (AI music composers as creative partners, ubiquitous real-time translators). Importantly, as these applications spread, they bring along considerations about how to use them responsibly, which we address next.

Ethical Considerations

The rapid advancement of AI in generating and interpreting audio raises several **ethical and societal concerns**. It is crucial to address these issues proactively to ensure these technologies are used for benefit and not harm. We discuss key ethical considerations, focusing on those pertinent to OpenAI's, Google's, and DeepSeek's audio AI:

- **Voice Cloning and Misinformation:** The ability to clone a person's voice from a small sample (demonstrated by OpenAI's Voice Engine and similar systems) poses a risk of **deepfake audio**. Malicious actors could impersonate someone's voice – for example, generating fake audio of a public figure saying something they never said, or tricking someone into thinking a family member called requesting emergency money (a known phone scam scenario). The realism of today's synthetic voices means humans often cannot distinguish fakes by ear alone. This challenges the trust we place in audio evidence. OpenAI and others are acutely aware of this risk. In response, OpenAI has not released Voice Engine widely and has built in safeguards: any usage requires the explicit consent of the person whose voice is cloned, and the **generated audio is watermarked** for detection. They also call for **"voice authentication experiences"** – essentially, if voice is used in an application, there should be verification steps to ensure the real person is behind it. Google similarly embeds identifiers (SynthID) in AI audio so that if an audio clip is in question, a detection algorithm can confirm it as AI-generated. Despite these measures, the broader concern remains: bad actors won't necessarily respect ethical guidelines or use watermarking. Society will need to adapt by being skeptical of audio recordings unless verified, similar to how deepfake videos have made us question video evidence. Regulatory responses are being discussed – for instance, laws that criminalize certain uses of voice deepfakes (especially if used to defraud or slander). The ethical imperative for AI developers is to keep improving detection, share tools with the public to check authenticity, and educate users that "hearing is not always believing" in the AI era.
- **Privacy and Surveillance:** Ubiquitous audio understanding can be a privacy double-edged sword. On one hand, users worry that devices or services are "always listening." Models like Whisper can run offline, which is better for privacy (no audio leaves your device). But cloud-based services (Google

Assistant, Alexa, etc.) do send audio to big tech companies for processing, raising concerns about data handling. If DeepSeek or others process personal conversations, there's the issue of how those transcripts might be used or stored. Companies must ensure strong privacy policies, limit data retention, and possibly allow opt-outs. Another angle is **surveillance**: in theory, a government could use mass AI listening to monitor citizens (e.g., identifying protest chants on streets from city CCTV mics, or flagging keywords in phone conversations). This is not new (wiretap and keyword-spotting tech existed), but AI makes it more scalable and accurate. Ethically, this use of AI is contentious – it pits security against civil liberties. Democratic oversight and transparency are needed to prevent abuse of such powerful listening capabilities. The companies themselves usually abide by privacy regulations (like not storing user voice data without consent, etc.), but it's a space to watch.

- **Bias and Fairness:** AI models can inherit biases present in their training data. For ASR, this often manifests as differing accuracy for speakers of different demographics. For example, if a model was trained mostly on American English, it might struggle with heavy accents or dialects (as was historically the case – systems performed worse on African American Vernacular English, etc.). Google's USM tried to cover many languages and dialects to reduce this bias, and OpenAI's huge training set for Whisper similarly included diverse accents. Still, bias can creep in. It's important to evaluate models on diverse datasets (male/female voices, various ethnic accents, different ages) to ensure equitable performance. For TTS, fairness might involve representing a wide range of voices – not just having default voices that reinforce stereotypes (e.g., always providing a female assistant voice by default, which some argue reinforces gender roles of women as assistants). With cloning, there's also a concern of misuse – e.g., could someone clone the voice of a minority figure to make fake statements that incite hatred? That crosses with misinformation too. Ethical development calls for including voices from different backgrounds in training and testing, and possibly giving users control (like the ability to choose the gender or style of assistant's voice that they are comfortable with). Encouraging diversity in AI-generated media (for instance, not all audiobook AIs should have the same standard accent) can also be positive.
- **Intellectual Property and Consent:** Audio data often contains content protected by copyright – songs, performances, etc. Training a model on someone's recorded voice or music potentially raises IP issues. OpenAI's Voice Engine was trained on "licensed and publicly available data", but details are not public. There is an ongoing debate: does training on copyrighted audio (without permission) count as fair use or infringement? Several lawsuits have been filed in related domains (e.g., Getty Images sued Stable Diffusion's makers for image training data). In music, artists are concerned that AI may be *learning their "style"* without compensation and could produce new songs "in the style of" a particular artist, which might compete with the artist's work. One article described DeepSeek's music model as a "wake-up call" because the music industry data is being leveraged by AI without clear frameworks. Ethically and legally, it will be important to establish norms – perhaps artists can opt out of their works being used in training sets, or they could receive royalties if an AI-generated piece clearly borrows from their creations. OpenAI's voluntary commitments hint at respect for creators (they allow artists to opt-out from image model training, for instance), and one could foresee similar moves for audio. When it comes to *personal voices*, consent is paramount. It is generally unethical (and in some jurisdictions illegal) to clone someone's voice without their permission, especially for public figures (impersonation). OpenAI's usage policies for Voice Engine prohibit impersonating individuals without consent ⁷ and even suggest a "no-go voice list" (e.g., disallowing the cloning of voices of prominent figures who haven't agreed to it). This could be implemented by checking any

sample audio uploaded by users against a database of protected voices and blocking matches. Such measures will be important to prevent misuse.

- **Impact on Human Labor and Creativity:** The deployment of AI audio technologies will inevitably affect certain jobs and creative processes. For example, transcriptionists and captioners might find less demand for their services as ASR automates their work. Voice actors and narrators are concerned that voice cloning could reduce opportunities – why hire a voice actor for a minor role when an AI can mimic a famous actor’s voice or generate a generic voice cheaply? Musicians worry that AI-generated music could flood the market or be used in place of human composers for background music, jingles, etc. These concerns are similar to those raised in other fields with generative AI (visual art, writing). Ethically, it’s important to ensure these technologies are used to assist and not purely displace humans, and to find new roles for human expertise (for instance, a voice actor might license their voice to an AI company and get paid for its usage, or musicians might use AI to enhance their productivity rather than replace it). We should also value transparency: if something is AI-generated, should it be labeled as such? Both OpenAI and Google have endorsed labeling AI-generated content (hence watermarks), which helps maintain honesty with audiences and preserves the value of human-made art when advertised as such. The **creative landscape** might shift to one where human and AI collaboration is the norm – ethically, we should navigate that transition in a way that respects the contributions of human creators and doesn’t deceitfully pass off AI work as human or *vice versa*.
- **Safety and Abuse Prevention:** Both OpenAI and Google emphasize safety research in the context of audio AI. Beyond misinformation, there are other misuse scenarios: harassing people with generated messages, creating obscene or hateful audio content, etc. While these companies cannot control everything users do with the tech, they can implement guardrails. OpenAI likely has filters to prevent Voice Engine from being used to generate certain content (for example, the API could reject requests that seem to be generating hate speech in someone’s voice). Google’s Gemini, by virtue of being an assistant, has conversation safeguards – it recognizes user tone and could refuse if someone tries to make it say something malicious in audio form. DeepSeek being open-source (in parts) could be used by anyone, which raises the usual open AI issue: how do you restrict misuse if the model is out in the wild? Some argue open models empower good and bad actors equally, so responsibility falls more on users or policy. Closed models allow companies to implement usage policies, but then those policies need to align with societal values. The partnership with policymakers is vital; indeed, OpenAI and others have made *voluntary commitments* to safety and external input. For audio AI, this might involve working with organizations on deepfake detection, contributing to standards (like audio watermarking formats), and staying ahead of emerging threats (for instance, if voice AI is used to generate ultrasonic sounds as a weapon, an edge-case but worth thinking about).

In summary, the ethical landscape for AI in audio is complex and evolving. The major developers are aware of these issues: OpenAI explicitly started public conversations on responsible deployment of synthetic voices, and Google has built in transparency features. It will require continuous vigilance, interdisciplinary collaboration (tech, law, ethics, media), and likely new laws or norms to ensure these powerful tools are used in ways that enrich society while minimizing harm. As we develop even more capable audio AI (e.g., models that can generate complete AV experiences or perfectly mimic anyone), the ethical stakes will only grow, making this an essential part of any discussion on audio AI advancements.

Limitations and Future Directions

Despite the impressive progress in AI for audio, current systems have several **limitations**. Recognizing these limitations is important for guiding future research and setting appropriate expectations. We also discuss promising directions for future improvements, some of which are already on the horizon.

Current Limitations:

- **Accuracy in Adverse Conditions:** Even the best ASR models can falter under tough real-world conditions – heavy background noise, multiple people talking over each other, strong accents/dialects that were underrepresented in training, or speech impaired by medical conditions. Whisper and Google’s ASR have made strides in noise robustness, but anyone who’s used auto-transcription in a crowded environment knows errors still occur (e.g., a live concert or a busy cafe yields messy transcriptions). Likewise, speaker diarization is not perfect: it might confuse speakers if their voices are similar or if many talk at once. Sound classification models might misidentify unusual noises. These models lack common sense understanding of context that a human listener might use (for example, knowing that at a beach, a certain muffled sound is likely wind or waves, not applause). Handling cocktail-party scenarios (separating voices) is still a challenge, though techniques like source separation (e.g., “unmixing” audio into individual speakers) are improving and might be integrated into future ASR.
- **Latency vs. Quality Trade-offs:** To achieve real-time performance, models often operate with small context windows or use approximations that can degrade output quality slightly. For instance, streaming ASR might have lower accuracy than batch ASR because it can’t leverage future context as much. In Google’s Gemini, the “Flash” mode sacrifices some fidelity for speed. So users face a trade-off: either slight delays for the highest quality (Gemini Pro, Whisper large, etc.) or instant results that might be slightly less accurate or less natural (Gemini Flash, Whisper small). There’s ongoing research to narrow this gap, but it’s a limitation that real-time and highest-quality are not always simultaneously achievable with current technology, especially on edge devices with limited power.
- **Generative Quality and Coherence:** For longer outputs, maintaining coherence is difficult. In TTS, generating a very long passage can lead to odd prosody or loss of the speaker’s style unless carefully controlled (some TTS systems break paragraphs into sentences to handle this). For music, models like MusicLM or DeepSeek might do well for, say, a 3-minute piece, but would they hold structure for a 10-minute symphony with evolving themes? Possibly not yet; they might regress into looping patterns or lose musical development. Jukebox had this issue, often meandering or creating surreal shifts beyond a certain duration. We should note that evaluation of generative audio quality is partly subjective; while MOS can measure naturalness, judging musical quality or appropriateness is nuanced. So, current models might produce “technically correct” audio that still lacks the emotional impact or creativity of human-produced content. For instance, AI music might sound bland or derivative at times, since it’s essentially interpolating seen examples rather than truly innovating (though some argue at sufficient complexity, the line blurs).
- **Limited Understanding of Semantics in Audio:** While speech AI is good with words, AI’s understanding of non-speech audio (like interpreting what a particular sound means in context) is still limited. For example, if an audio has the sound of glass breaking and a dog barking, a human inferential leap might be “perhaps a window broke and it scared the dog.” An AI might detect “sound

of glass break” and “dog bark” as separate labels, but not piece together a narrative. There are efforts in multimodal models to handle this (e.g., image+audio understanding), but it’s early. DeepSeek or GPT-4 with plugins could be tasked with reasoning about an audio scene transcript or description, but a lot of high-level understanding is left to humans. This is a limitation when trying to use AI for tasks like *audio surveillance* or *content analysis* in videos – it might miss the forest for the trees.

- **Energy and Compute Consumption:** The largest models, like DeepSeek’s, consume enormous computational resources for training and inference. This has environmental and economic implications. Whisper can run on smaller devices, but training Whisper took a lot of GPU time on those 680k hours. Google’s 12M hours of training for USM is huge – albeit self-supervised, which is efficient per unit data but overall massive. The energy cost of running these models at scale (e.g., millions of voice assistant queries per day) is non-negligible. Latency improvements often come by throwing more compute (like running on big clusters to parallelize response generation). For sustainable AI, there’s a need to make models more efficient – through model compression, better algorithms, or hardware acceleration (like more efficient ASICs for audio processing).
- **User Acceptance and Trust:** This is more of a soft limitation – users may not fully trust or accept AI in certain roles yet. For example, an AI-generated news report read by an artificial voice might be met with skepticism, or a musician might be reluctant to use an AI-generated riff in their song due to authenticity concerns. Some users simply prefer the human touch (e.g., many still enjoy audiobooks narrated by humans because of the acting performance, something AI voices, while improving, may not fully replicate the subtle artistry of a trained actor). Thus, even if the tech is capable, adoption might lag due to human factors, and we must consider how to earn trust (through transparency, gradual improvements, or hybrid approaches where humans oversee or fine-tune AI outputs).

Future Directions:

- **Advanced Multimodal Integration:** One exciting direction is tighter integration of audio AI with other modalities. Google Gemini already points this way: imagine an AI that can watch a video (visual input), listen to its audio, and understand the whole scene holistically – that’s very useful for tasks like video captioning or guiding a robot. For instance, future models could take in both sound and vision to do **audio-visual scene understanding** (e.g., see a door closing and hear a slam, then deduce someone left the room angrily). OpenAI’s GPT-4 has vision input; perhaps an audio input modality might come (we could foresee a “GPT-5” that you can speak to directly, and it hears the tone of your voice, etc.). DeepSeek too would benefit from better sensory fusion – maybe having a module that directly processes audio waveforms and interacts with its text modules. This multimodal fusion would lead to richer AI assistants (ones that can, say, take a phone call and also look up relevant info while listening).
- **Personalization and Adaptability:** We anticipate models that can personalize on the fly to individual users. For ASR, that might mean learning a user’s unique accent or vocabulary quickly (imagine a personal ASR model that adapts to your voice after a few uses and gets significantly more accurate for you than a generic model). For TTS, personalization could mean generating speech that matches not just a given speaker but also a user’s preferred style (maybe you like your AI to talk faster or to use a certain playful tone). Both OpenAI and Google have taken steps here – e.g., style controls in Gemini TTS, and OpenAI’s focus on cloning specific voices. Future improvements might

allow *everyone* to have their own custom AI voice: you speak a few lines and your devices thereafter *talk back in your voice* (some might find that eerie, but others could use it for continuity, say a GPS navigation in your own voice). On the understanding side, adaptability also means **contextual learning**: if an AI is used in a workplace, feeding it domain-specific audio (like meetings of a company) should help it learn jargon and speaker roles to better transcribe and summarize future meetings. Techniques like on-the-fly learning, federated learning on device, or efficient fine-tuning (like Low-Rank Adaptation, LoRA, applied to audio models) could facilitate this personalization without requiring full retraining.

- **Improving Naturalness and Emotion:** While current synthetic voices are very natural, there's always room to be more human-like. One frontier is emotional expressiveness – not just adding a happy or sad tone on command, but dynamically modulating emotion as a human would. This involves understanding context deeply (so the AI knows *when* to sound excited vs. serious). Prosody prediction is still an area of research; the nuances of stress and intonation sometimes escape neural TTS, leading to speech that is good but subtly “flat” or off in emphasis. Future models might incorporate better linguistic and world knowledge to inform prosody (almost like an actor who first understands the meaning of a line and the subtext before reading it). Also, **laughter, sighs, hesitations** – these paralinguistic features are part of natural speech. We may see TTS that can emit a polite chuckle or a thoughtful “hmm” where appropriate (some current systems can insert a breath sound or a sigh from a fixed library, but making it contextually appropriate is key). The Gemini article mentions non-speech vocalizations like laughter as something it recognizes and presumably could emulate. In generated music, increasing the *creativity* and *originality* of compositions is a goal – avoiding plagiarism of training data melodies and perhaps introducing controlled randomness or “inspiration” to create novel art. There is interesting work in AI that attempts to ensure generated content is not just a regurgitation but has novelty; applying that to music (maybe by combining influences in new ways) could produce truly interesting compositions.
- **Robustness and Verification:** Future research will also focus on making models more robust and secure. By robust, we mean resistant to adversarial attacks or weird inputs. For instance, adversarial noise could trick an ASR into hearing words that aren't there; making models robust to that is important especially for security domains. Verification and detection tools will improve – perhaps every AI audio will come with a cryptographic signature or hidden mark so that any platform (Twitter, YouTube) can automatically flag “this audio was AI-generated” before it spreads. Watermarking research is ongoing and making progress in making marks that survive compression and mild edits. As audio deepfakes get better, detection must too; future detection might involve cross-modal checks (e.g., an AI that listens to a voice and looks at a supposed video of the person speaking and sees mismatches in micro-expressions vs. voice, etc.). This cat-and-mouse game will continue, and involvement from organizations like Adobe (with their Content Authenticity Initiative) or coalitions of AI firms could lead to standards for authenticity metadata in media files.
- **Hardware Acceleration and Edge AI:** On the deployment side, the future will likely bring more specialized chips for running AI audio tasks efficiently, enabling more on-device processing. Already, some smartphones have “neural processing units” that can handle tasks like voice wake-word detection with negligible battery impact. We may see high-end earbuds that do on-device noise suppression and even live translation between two people in conversation without sending data to the cloud. Tiny ASR models might live in appliances for offline voice control (privacy-first smart home devices). There is also interest in **low-power music generation** – imagine a portable device that can

compose songs for you on the go. To get there, models must be optimized or new algorithms discovered that are far lighter. Techniques like knowledge distillation, quantization (running models in 8-bit or 4-bit precision), and algorithmic shortcuts (like using RNNs or S4 models for long audio sequences instead of full attention) will be key research directions. DeepSeek’s MoE itself is a form of efficiency (not every expert is active per token), so perhaps more models will adopt sparsity to keep scaling without linear cost.

- **New Modalities and Capabilities:** Future AI might handle audio in ways we haven’t fully explored – like **audio-to-image** (creating images from sound, e.g., “draw what this soundscape sounds like”) or **audio-based reasoning** (using clues in a speaker’s tone to gauge their mood or truthfulness – though fraught with uncertainty). Also, integrating **brain-computer interfaces**: some research maps neural signals to synthetic speech (giving voice to the voiceless via brain signals). AI will play a role in that decoding and encoding process, potentially using models like Voice Engine to speak out what someone is attempting to say via neural cues. This crosses into neuroscience and is probably further out, but not science fiction given early experiments. In music, AI might help with mastering and audio engineering tasks (automatically mixing tracks, adjusting equalization to professional quality). These expanded uses will further blur the lines between creative tool and creative agent.

In summary, the field of AI audio is vibrant with opportunities for improvement. Many of the current limitations (noise robustness, data efficiency, expressiveness) are active areas of research and engineering. The future points toward AI systems that are more **integrated, personalized, efficient, and trustworthy**. We can expect to converse with AI more naturally, hear AI-generated content that moves and entertains us, and simultaneously rely on AI to keep the use of audio data honest and secure. Achieving all this will require collaboration between researchers, industry, and users to guide these models in a direction that augments human capabilities and experiences while mitigating downsides.

Conclusion

Artificial intelligence has rapidly advanced the frontier of what machines can hear and create in the realm of audio. In this paper, we examined how AI models have achieved impressive levels of proficiency in **audio understanding** (like transcribing speech, identifying speakers, and classifying sounds) and in **audio generation** (synthesizing realistic speech and even music). We focused on the contributions of OpenAI, Google’s Gemini, and DeepSeek – three leading efforts that encapsulate the trends in this field. OpenAI’s Whisper ASR and Voice Engine demonstrate the power of specialized models trained on massive data to achieve robust, near-human performance in transcription and voice synthesis. Google’s Gemini represents the integration of those capabilities into a single multimodal AI system that can engage in real-time audio dialog and controlled speech generation with remarkable quality. DeepSeek’s enormous mixture-of-experts models illustrate a strategy of sheer scale and unified learning, yielding breakthroughs like high-fidelity AI music generation ¹ and multimodal understanding across text and audio within one system.

Our comparative analysis shed light on the strengths and trade-offs of each approach. OpenAI’s models, being open or in preview, have democratized access to cutting-edge audio AI, but OpenAI has also navigated the **ethical deployment** challenges carefully (for instance, delaying public release of voice cloning until safety is assured) ⁷. Google’s approach, backed by its vast compute and data, currently leads in multilingual and real-time capabilities, making its AI attractive for interactive and global applications. DeepSeek’s approach pushes the envelope in model size and multimodality, hinting at a future where a single AI can seamlessly handle speech, text, images, and more – though with present limitations in speed

and access. In terms of **performance**, all platforms achieve high marks, with differences in specific areas (e.g., Google's ASR edging out others in many languages, OpenAI's voice cloning being uniquely advanced, DeepSeek's music generation potentially leading a new wave in creative AI). For **latency and deployment**, Google and OpenAI's cloud implementations enable real-time voice AI for wide audiences, whereas offline use at the edge is currently better supported by smaller models like Whisper or distilled DeepSeek variants.

We explored how these technologies are being applied – from voice assistants that allow natural spoken conversations with AI, to automatic captioning that makes media more accessible, to creative tools that can compose music or generate voice-overs at the click of a button. The **applications** are broad and growing, spanning industries from entertainment to education to healthcare. Each application brings value – whether it's saving time (through transcription), breaking language barriers (through translation and cross-lingual voice cloning), or providing assistance and accessibility (through screen readers or voice for the speech-impaired) ⁶. Alongside the benefits, we have critically examined the **ethical** dimension. The same tools that empower users can be misused for deception or violation of privacy. We emphasized that developers like OpenAI and Google are actively working on mitigations – such as watermarking AI-generated audio, usage policies restricting abuse ⁷, and public awareness initiatives – but the vigilance must continue. Policymakers, the AI community, and society at large will need to collaborate to set guidelines and norms (for example, perhaps requiring disclosure when audio is AI-generated in media, or making it illegal to clone someone's voice without consent).

Looking forward, the trajectory of audio AI points to even more **immersive and intelligent auditory experiences**. We foresee AI that can truly understand context and emotion in voices, making interactions more empathetic, and AI that can generate rich audio-visual narratives, blurring the line between listener and participant. Future models will likely be more *efficient*, addressing the current dependency on enormous datasets and compute through algorithmic innovations so that high-quality audio AI can run locally and securely. They will also likely be more *collaborative*, serving as creative partners to musicians, translators for diplomats, tutors for students, and assistive friends to those in need.

In conclusion, the field of AI audio has entered an exciting era where long-standing challenges like real-time speech recognition and natural speech synthesis are effectively solved, and new frontiers like open-domain audio dialog and AI-driven music creation are emerging. OpenAI, Google's Gemini, and DeepSeek exemplify the diverse strategies to advancing this field – from open models enabling widespread adoption to integrated AI assistants redefining user interactions, to massive models pushing performance boundaries. Each contributes to an ecosystem of tools that, if developed and deployed responsibly, can significantly enhance human capabilities and connectivity. As we continue to refine these technologies and address their societal impacts, we move closer to a world where speaking to an AI is as effortless as speaking to a person, and where the creative and communicative potential of audio is amplified by artificial intelligence in a multitude of positive ways.

References

1. Bapna, A., & Sainath, T. (2025, June 3). *Advanced audio dialog and generation with Gemini 2.5* [Blog post]. Google Keyword.
2. Ferguson, M. (2024). *DeepSeek's V3 AI music model: A symphony of disruption in the music industry!* OpenTools AI News. ¹

3. Google Research. (2023, March 6). *Universal Speech Model (USM): State-of-the-art speech AI for 100+ languages* [Blog post].
 4. OpenAI. (2022, September 21). *Introducing Whisper* [Blog post]. OpenAI.
 5. OpenAI. (2024, June 18). *Navigating the challenges and opportunities of synthetic voices* [Blog post]. OpenAI.
 6. Wiggers, K. (2024, March 29). *OpenAI built a voice cloning tool, but you can't use it... yet*. TechCrunch.
-

1 DeepSeek's V3 AI Music Model: A Symphony of Disruption in the Music Industry! | AI News

<https://opentools.ai/news/deepseeks-v3-ai-music-model-a-symphony-of-disruption-in-the-music-industry>

2 Speech-to-Text release notes | Cloud Speech-to-Text Documentation | Google Cloud

<https://cloud.google.com/speech-to-text/docs/release-notes>

3 4 5 MusicLM

<https://google-research.github.io/seanet/musiclm/examples/>

6 7 Navigating the challenges and opportunities of synthetic voices | OpenAI

<https://openai.com/index/navigating-the-challenges-and-opportunities-of-synthetic-voices/>